Concurrent CUDA Streams

A **stream** is a series of instructions, and CUDA has a **default stream**

By default, CUDA kernels run in the **default stream**

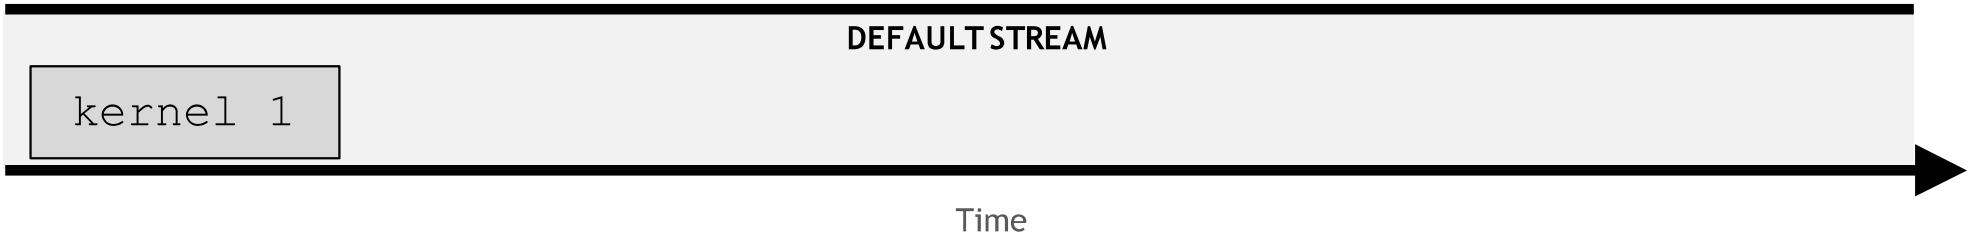

In any stream, including the default, an instruction in it (here a kernel launch) must complete before the next can begin

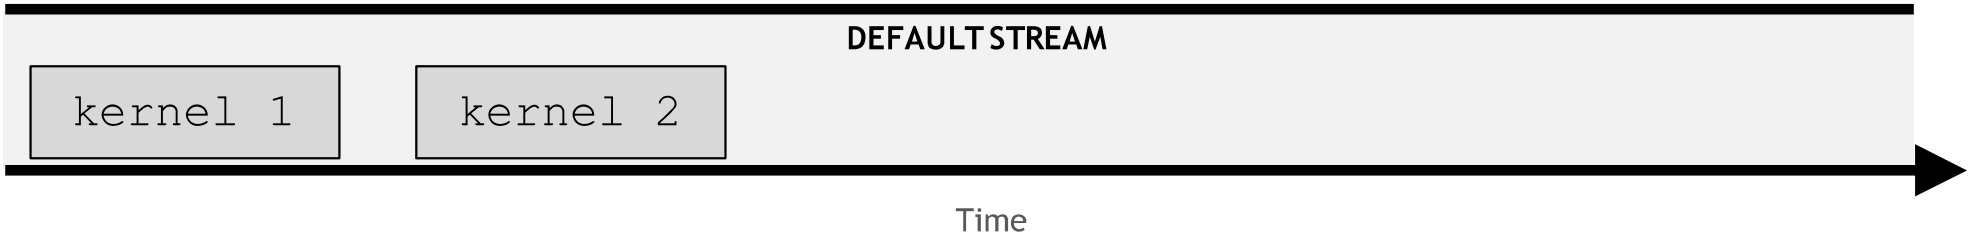

In any stream, including the default, an instruction in it (here a kernel launch) must complete before the next can begin

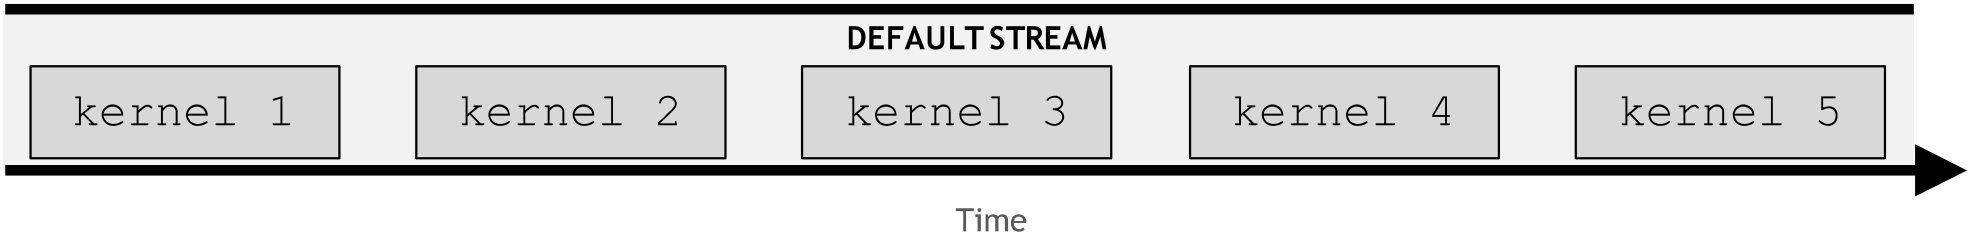

Non-default streams can also be created for kernel execution

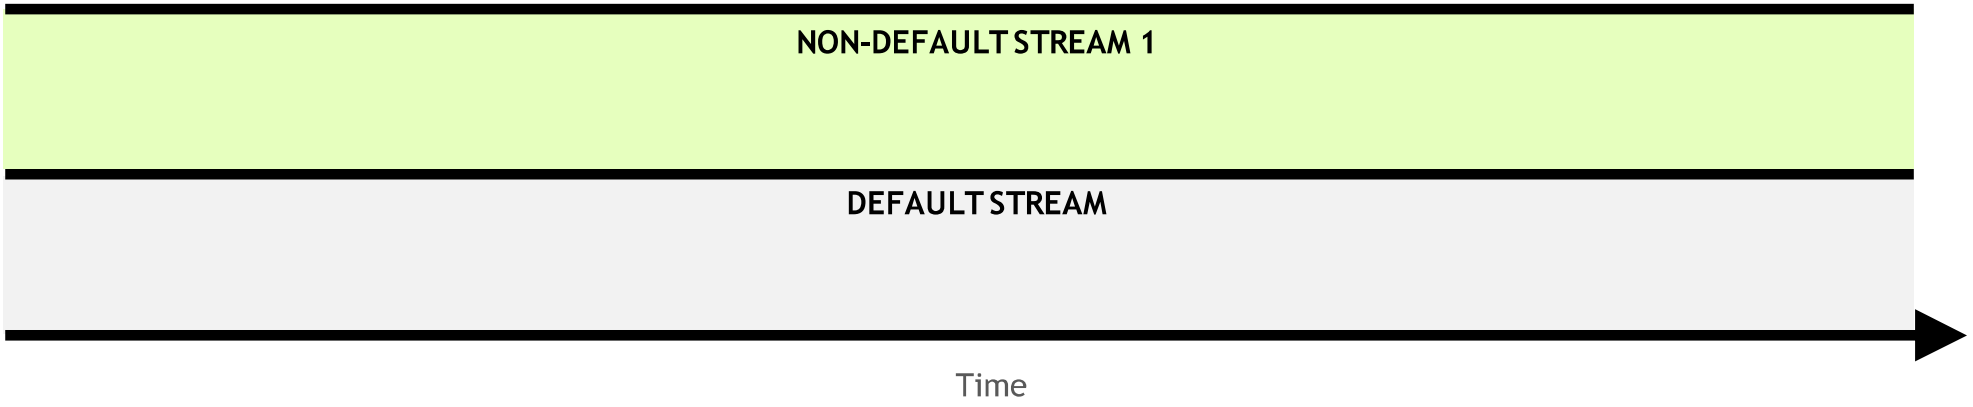

Non-default streams can also be created for kernel execution

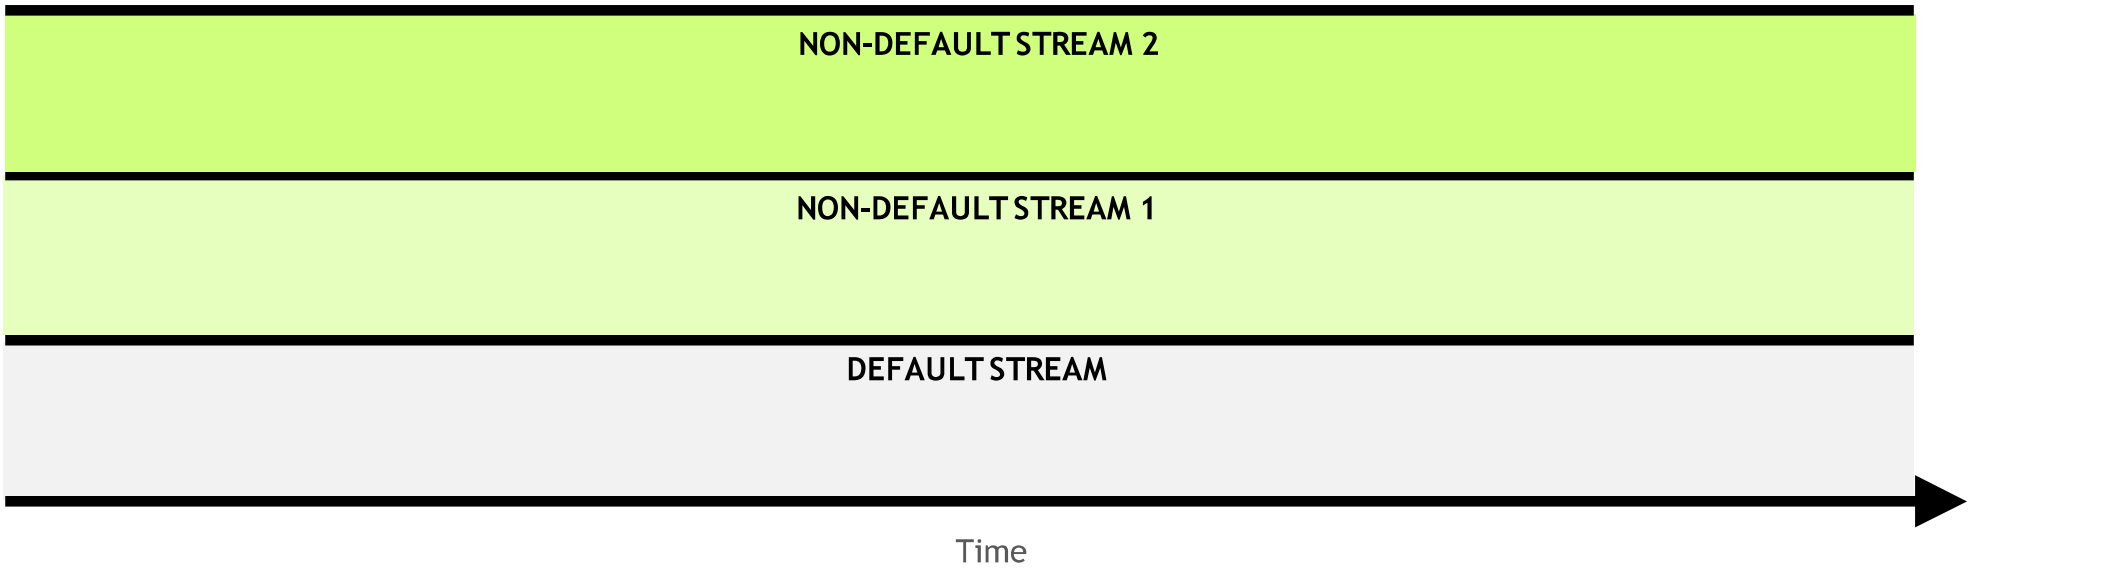

Kernels within any single stream must execute in order

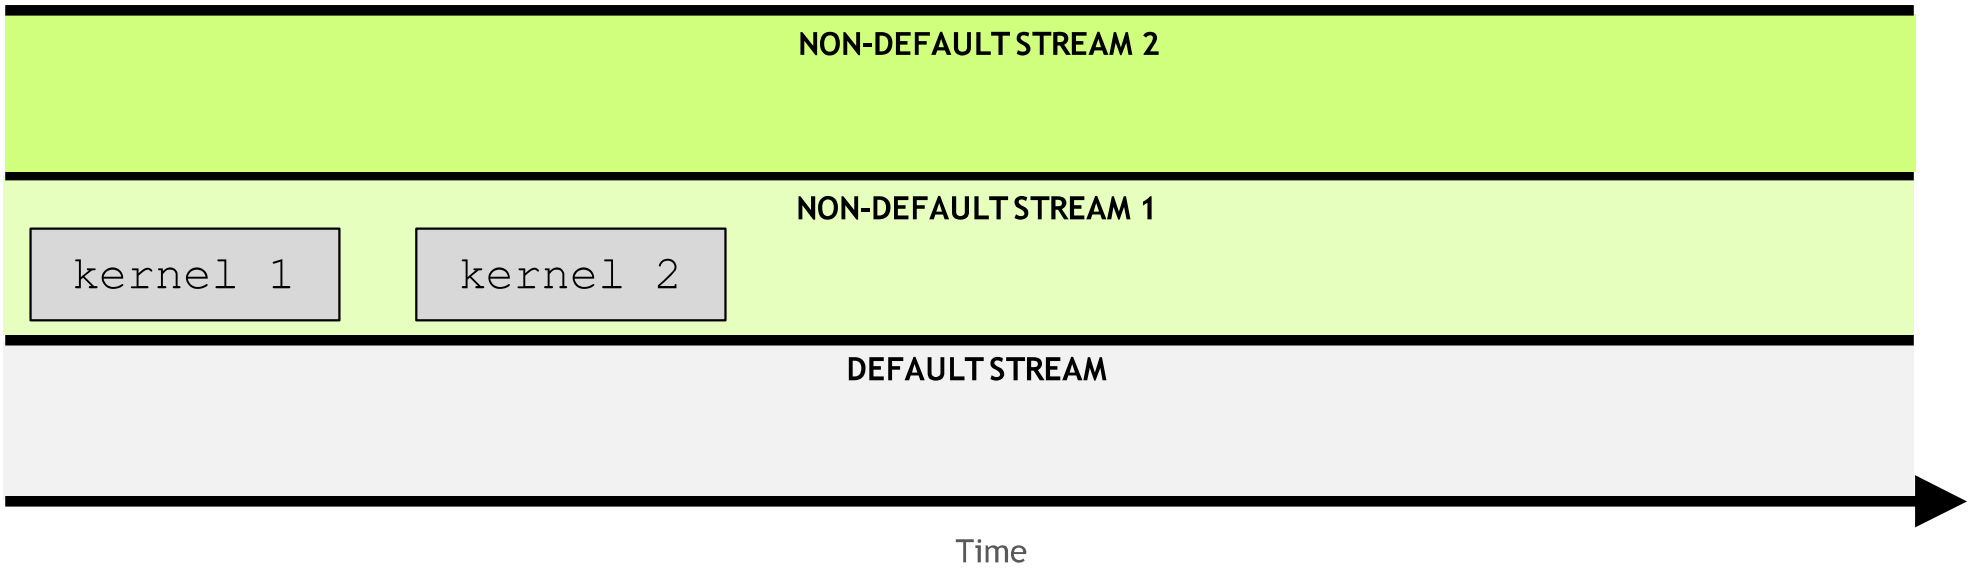

However, kernels in **different, non-default streams**, can interact concurrently

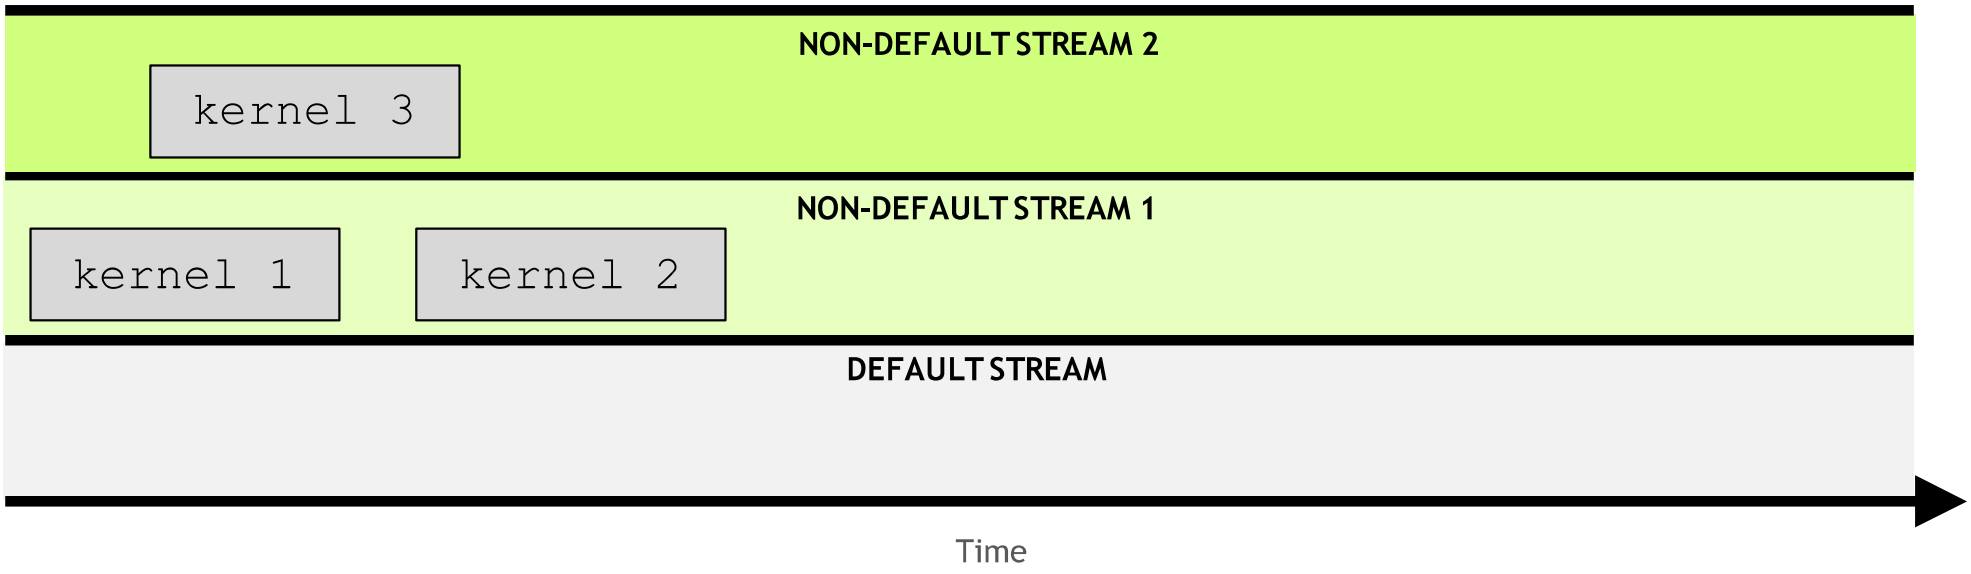

However, kernels in **different, non-default streams**, can interact concurrently

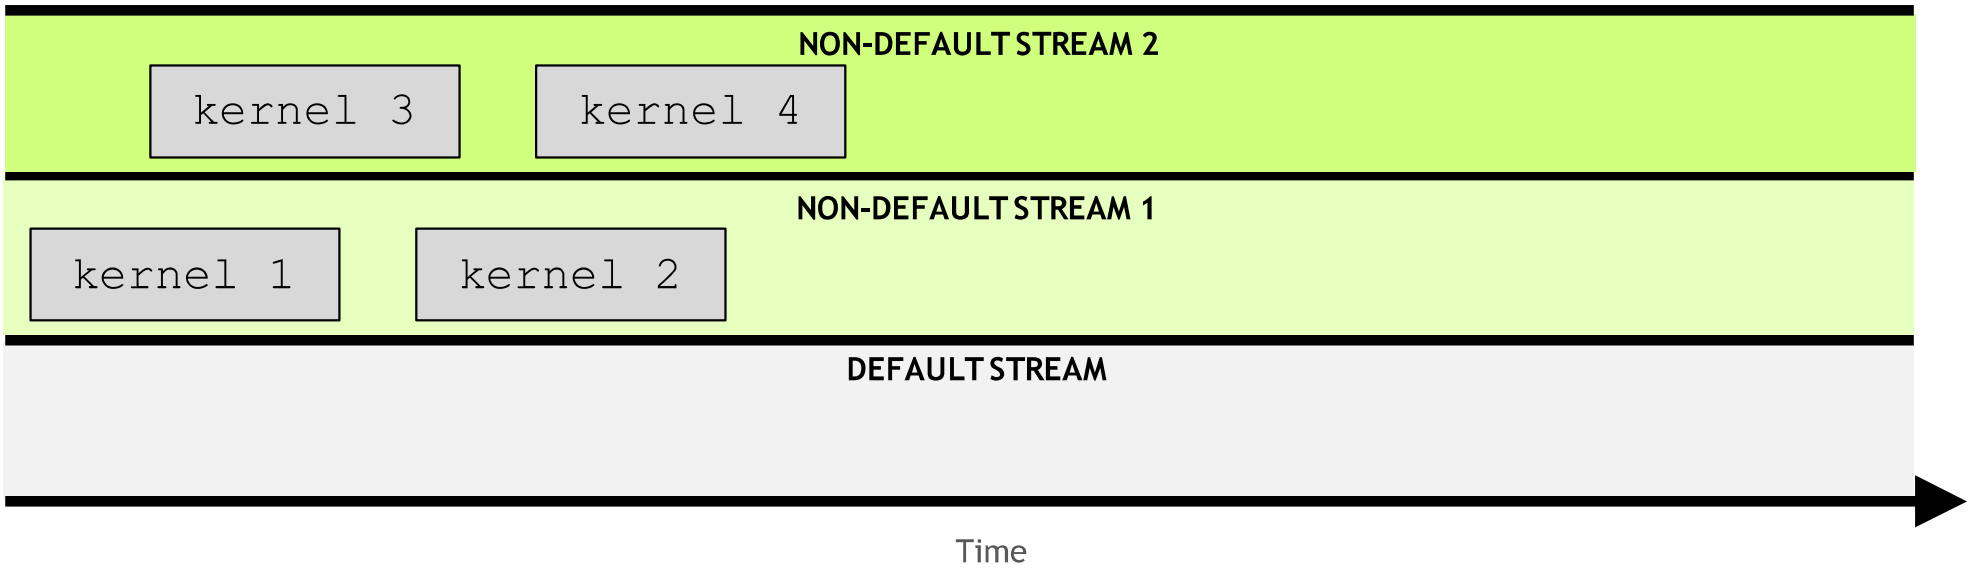

However, kernels in **different, non-default streams**, can interact concurrently

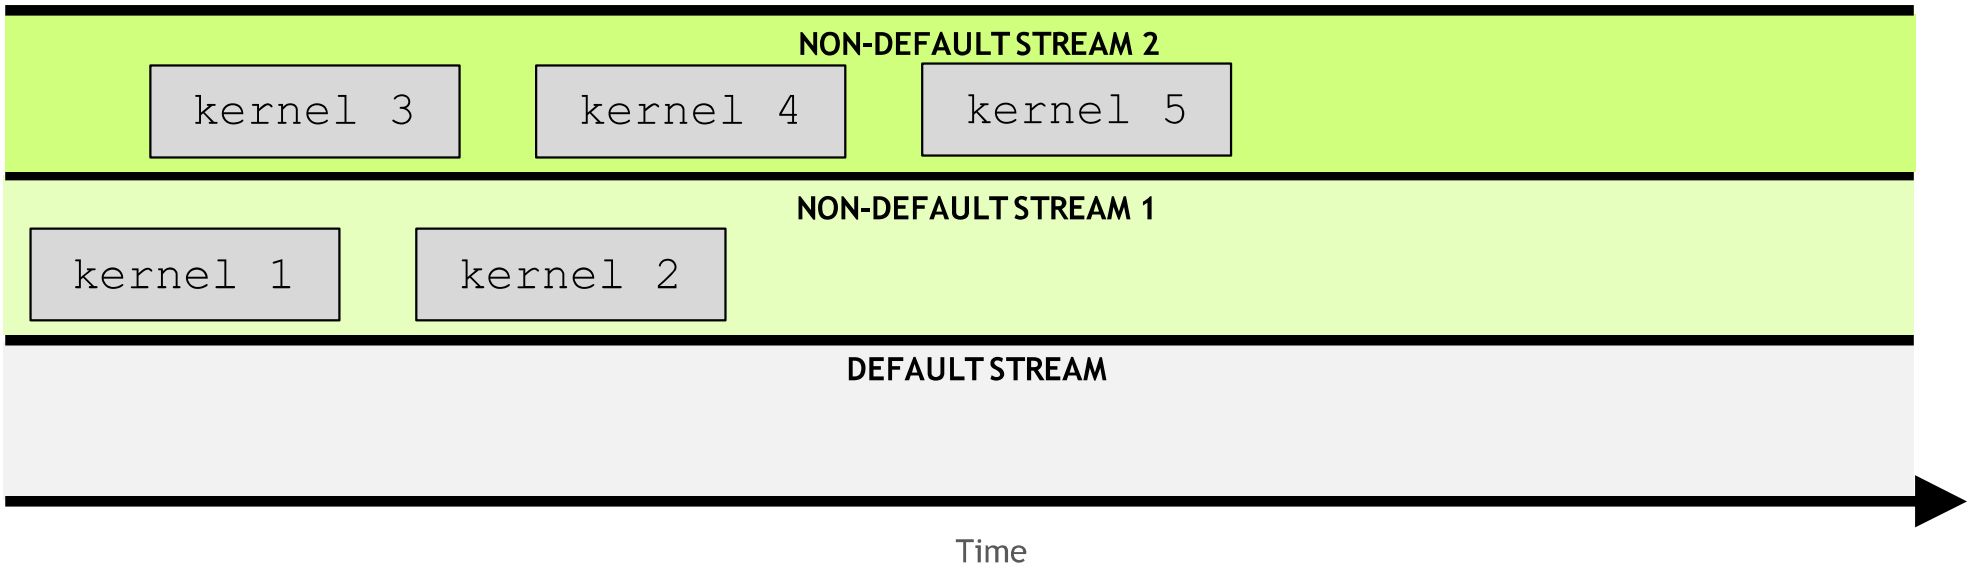

However, kernels in **different, non-default streams**, can interact concurrently

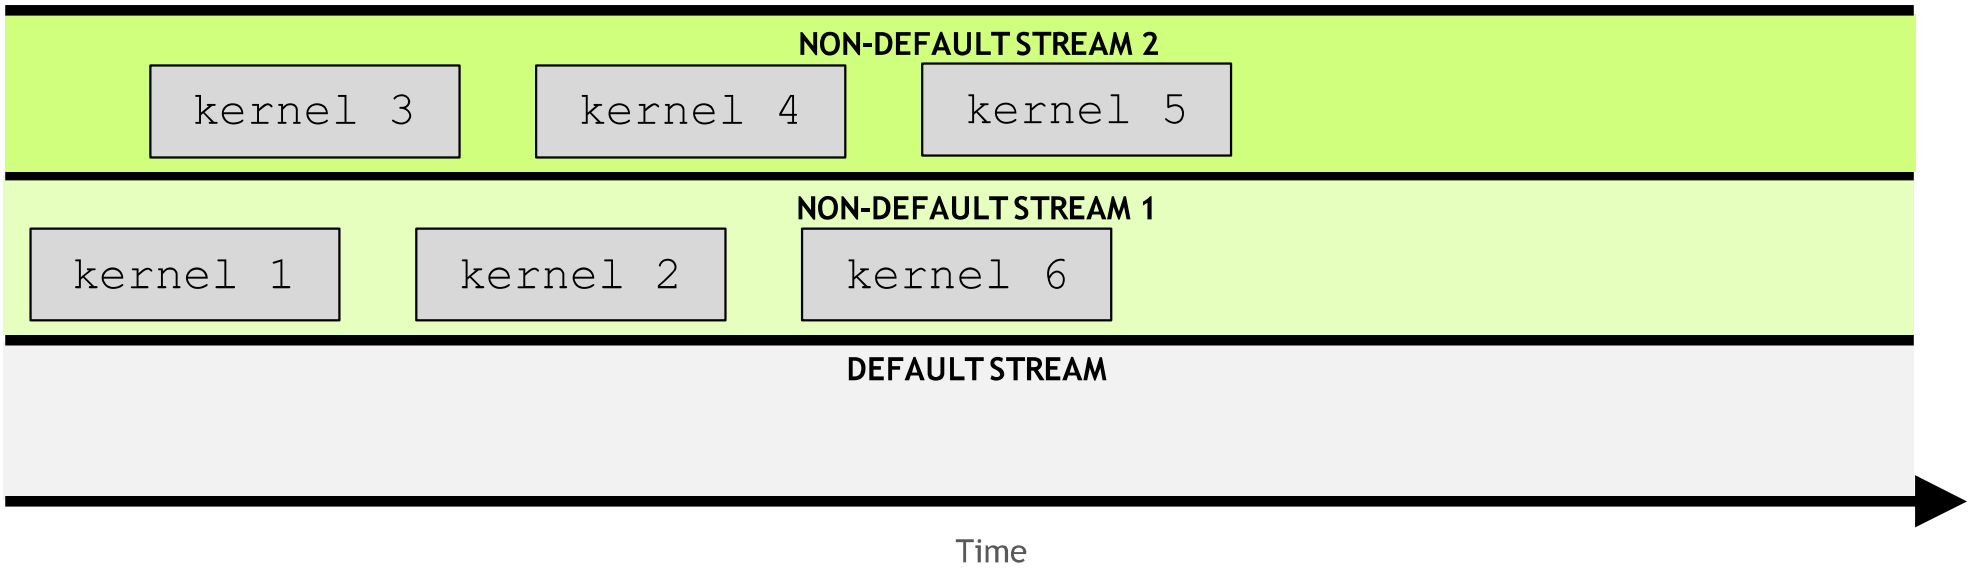

The default stream is special: **it blocks all kernels in all other streams**

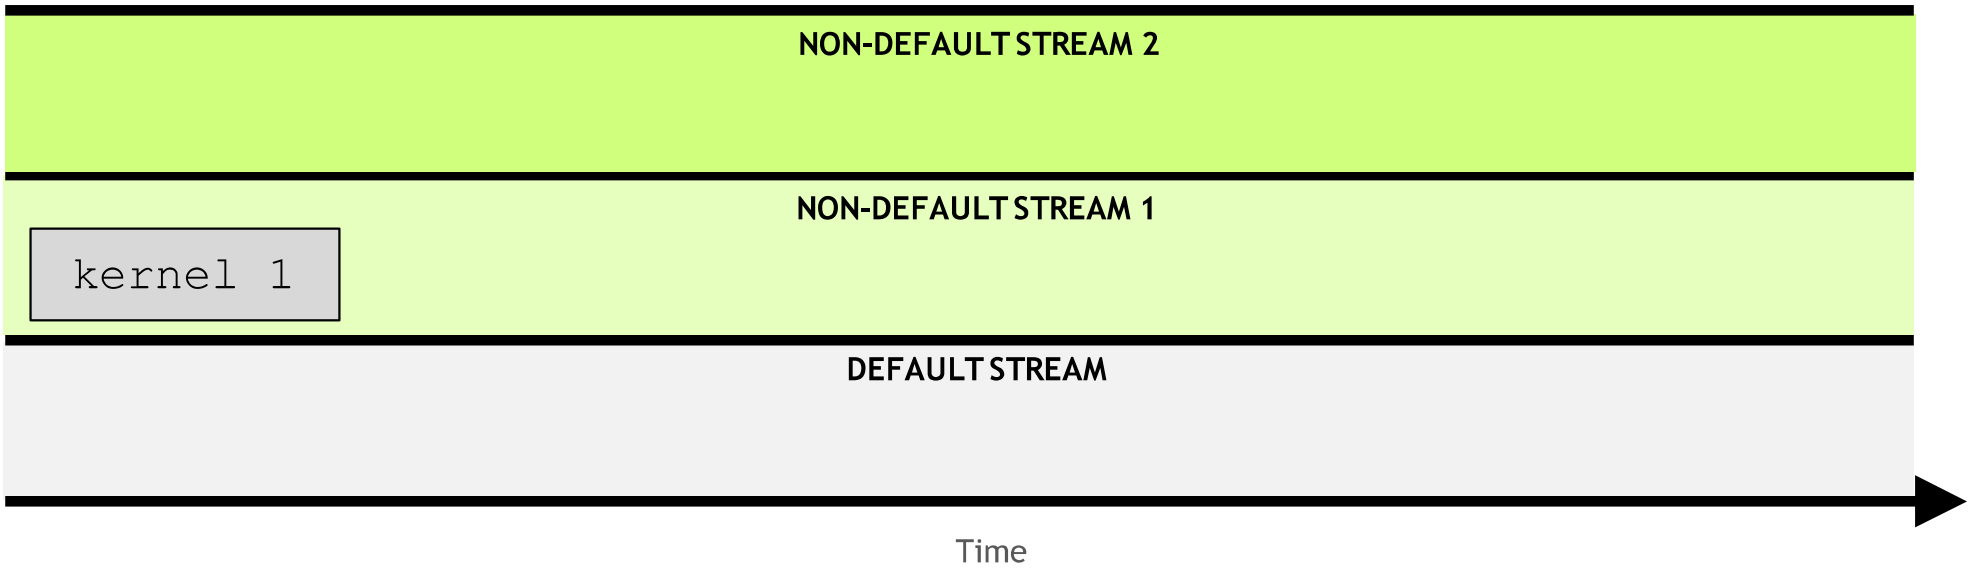

The default stream is special: **it blocks all kernels in all other streams**

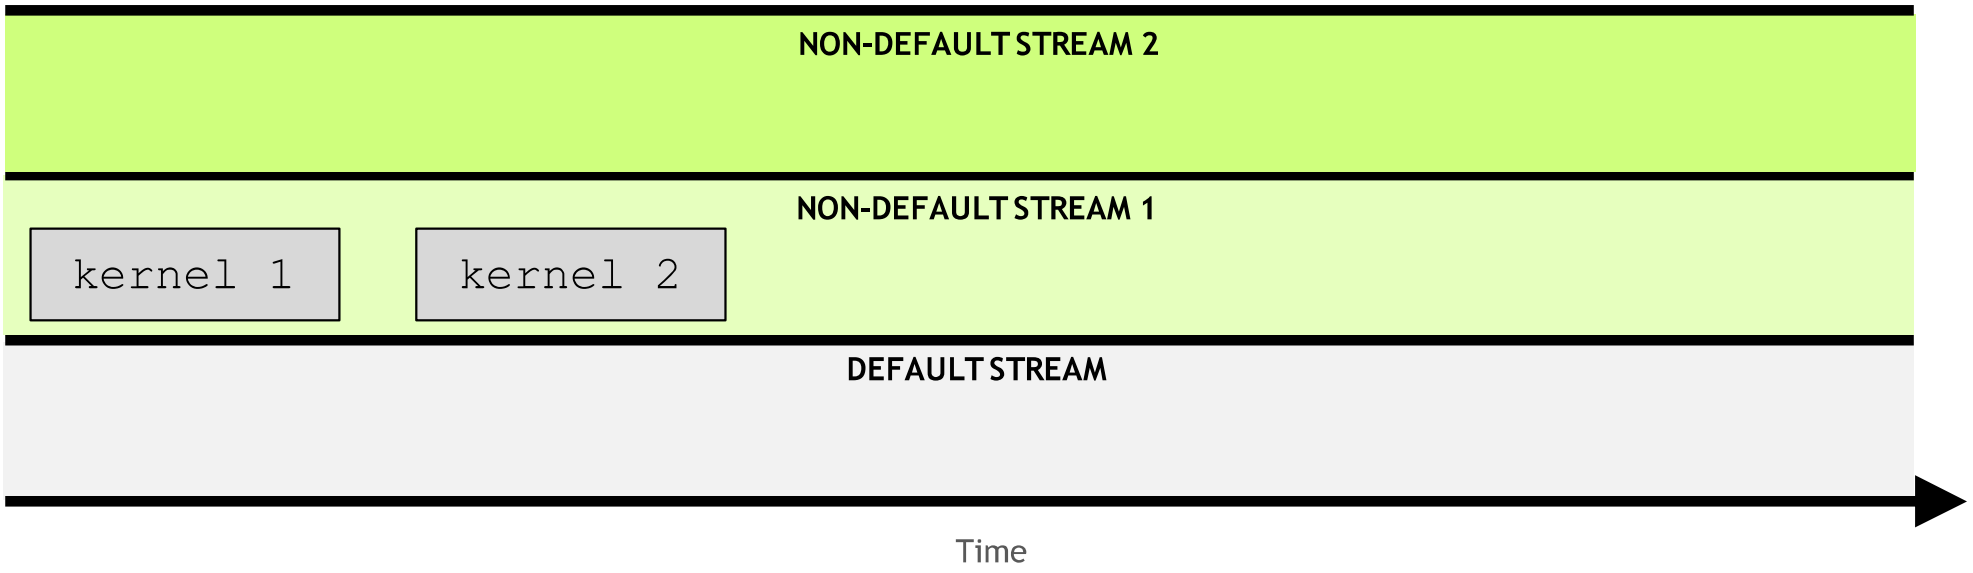

The default stream is special: **it blocks all kernels in all other streams**

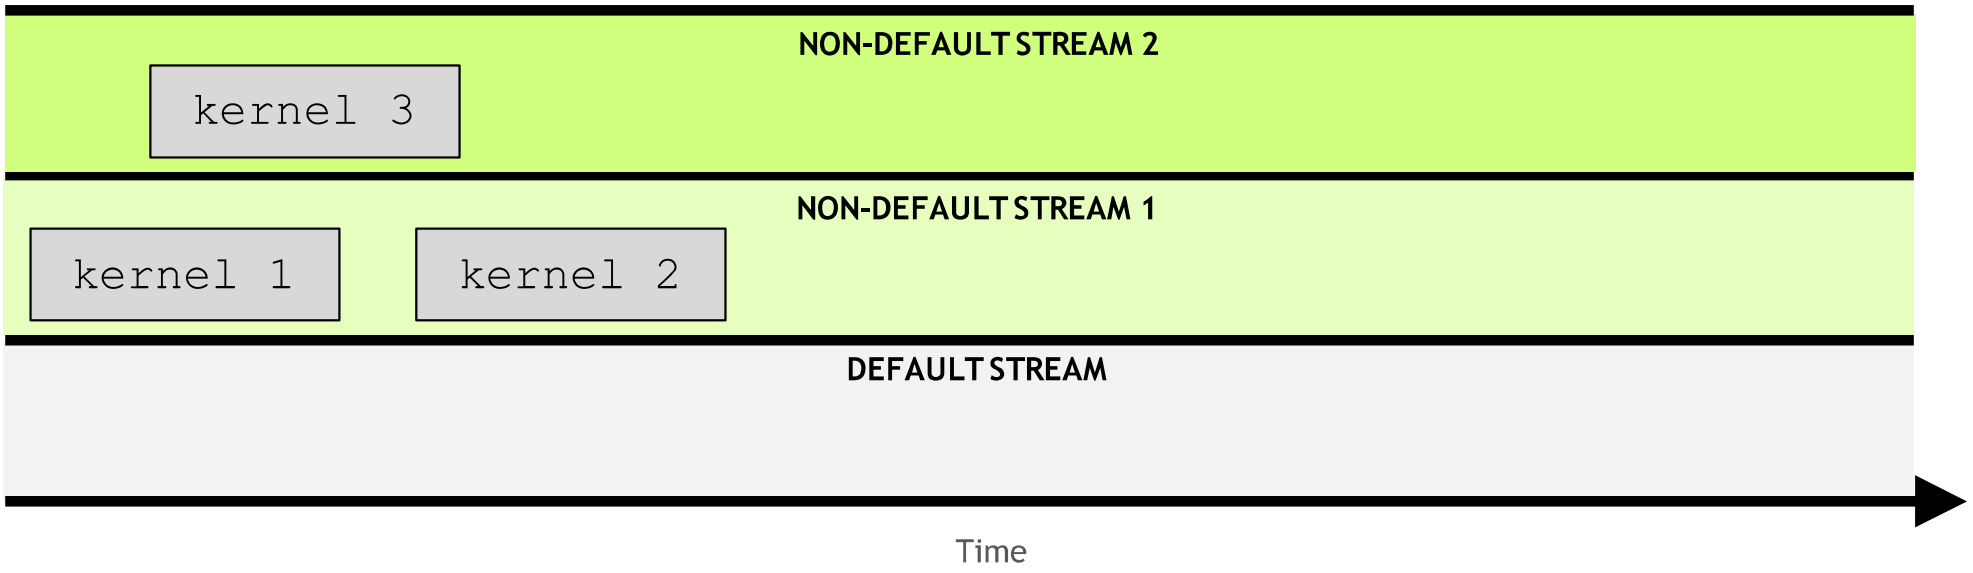

The default stream is special: **it blocks all kernels in all other streams**

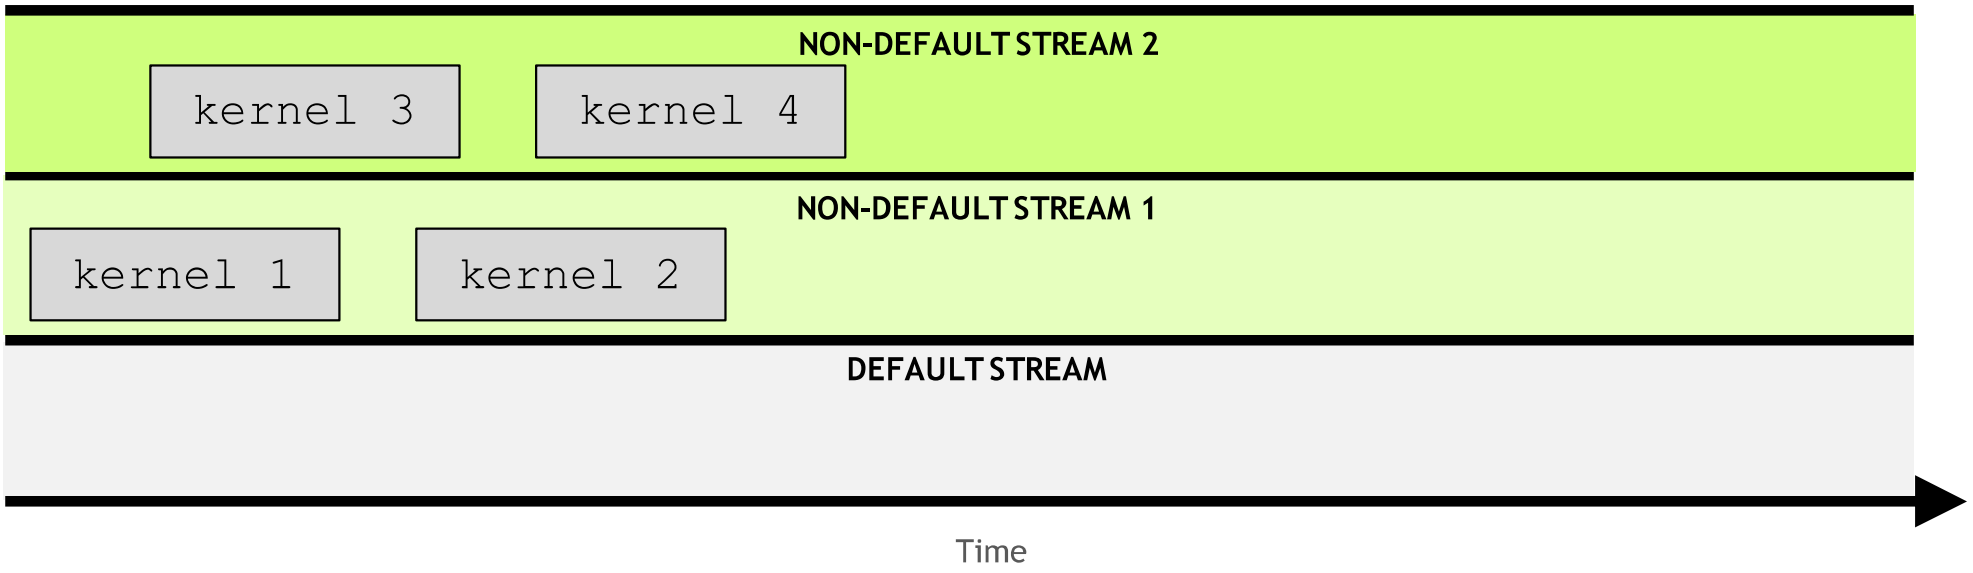

The default stream is special: **it blocks all kernels in all other streams**

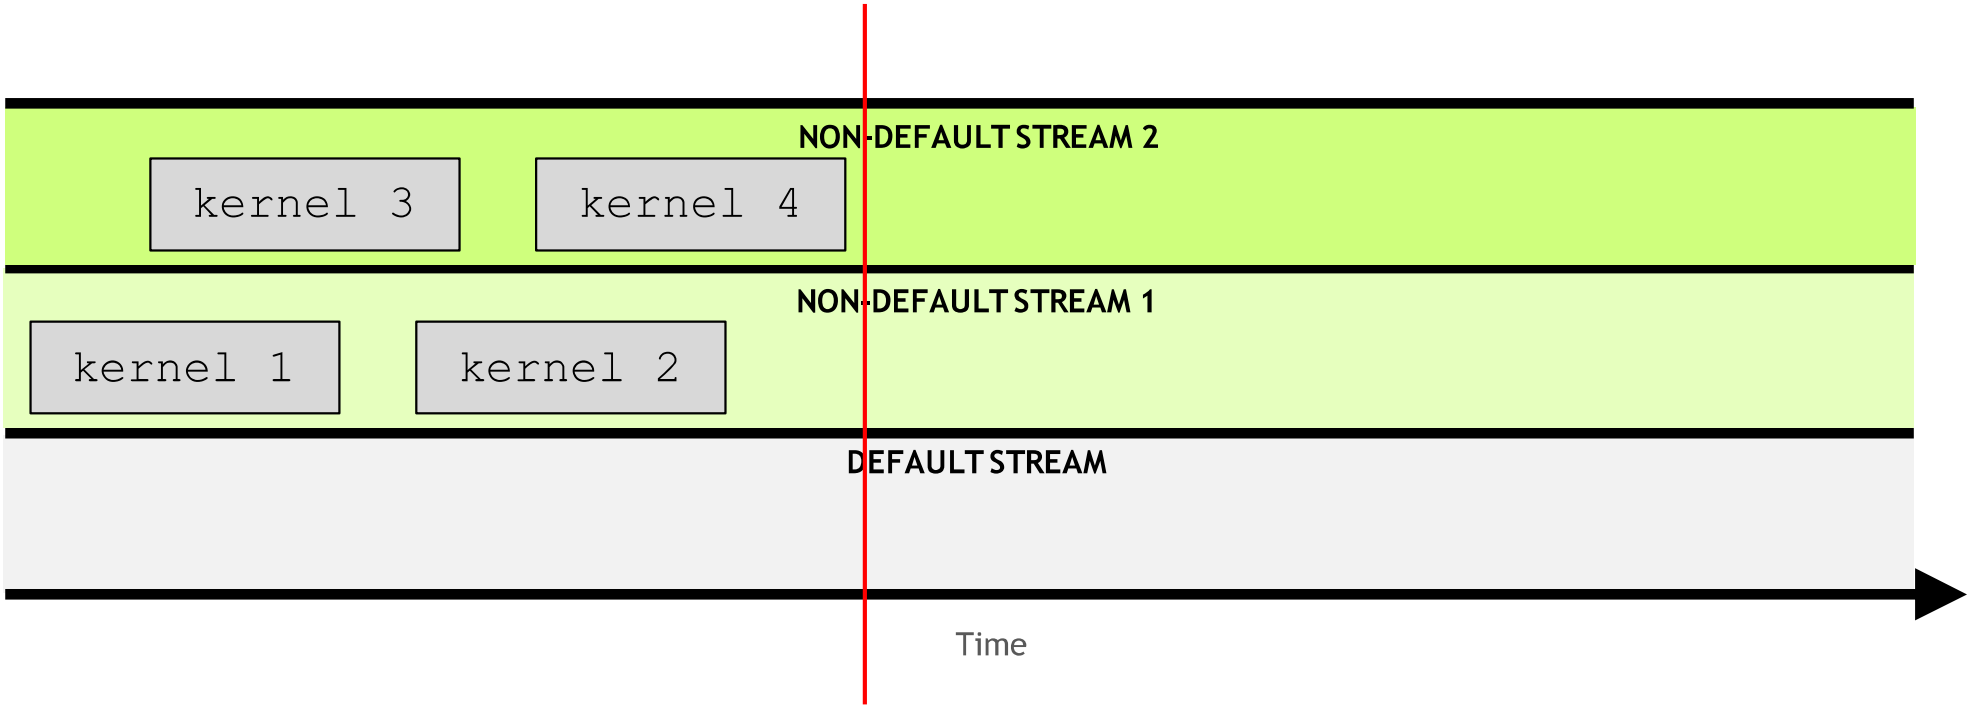

The default stream is special: **it blocks all kernels in all other streams**

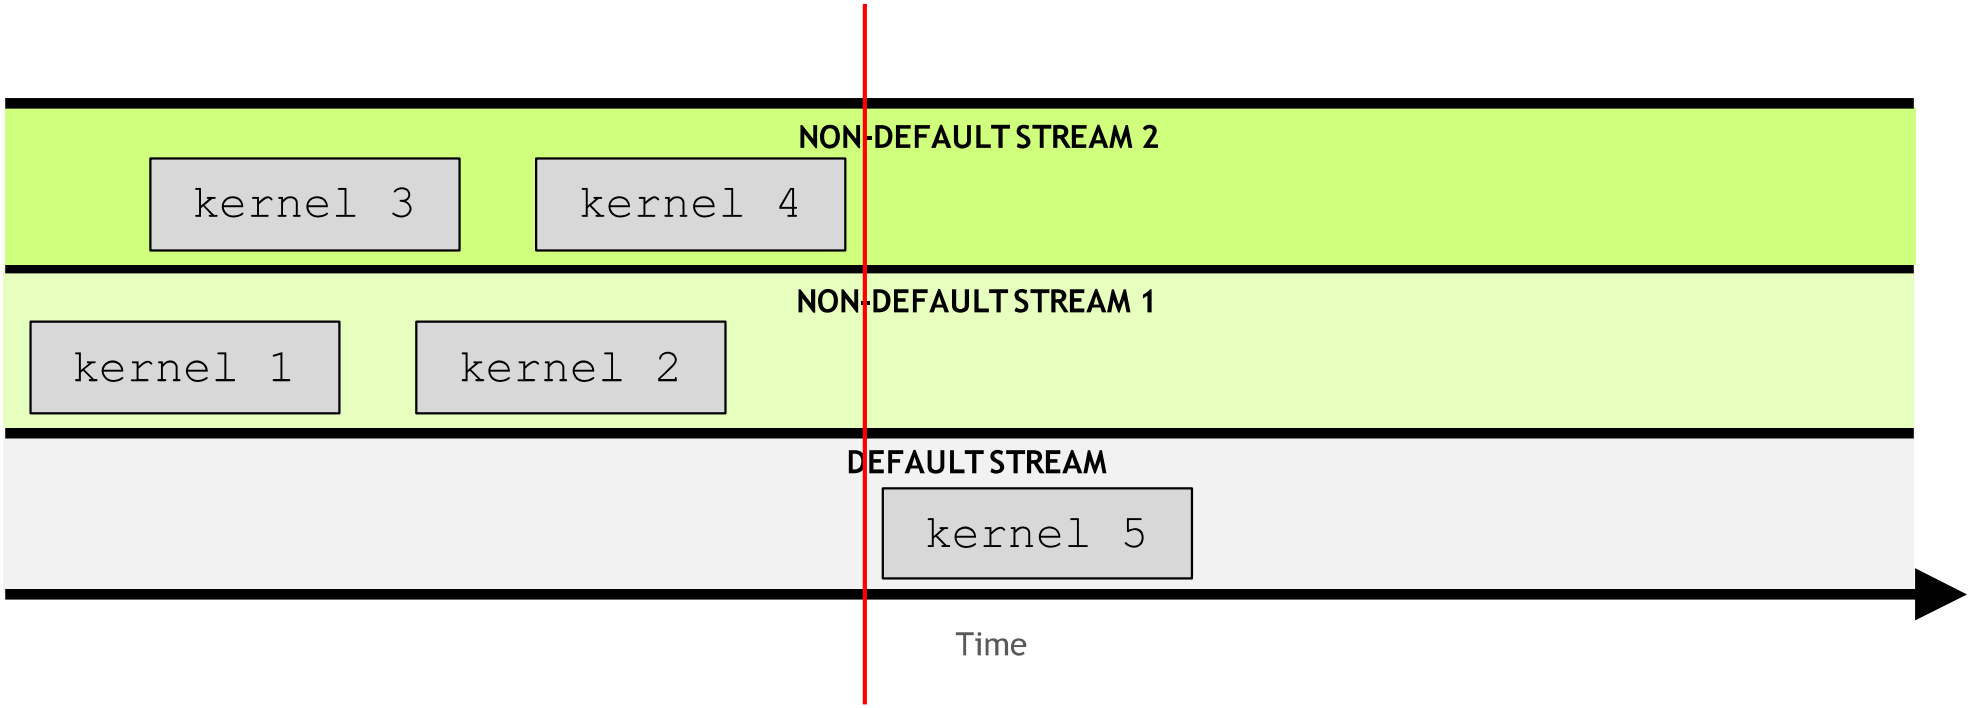

The default stream is special: **it blocks all kernels in all other streams**

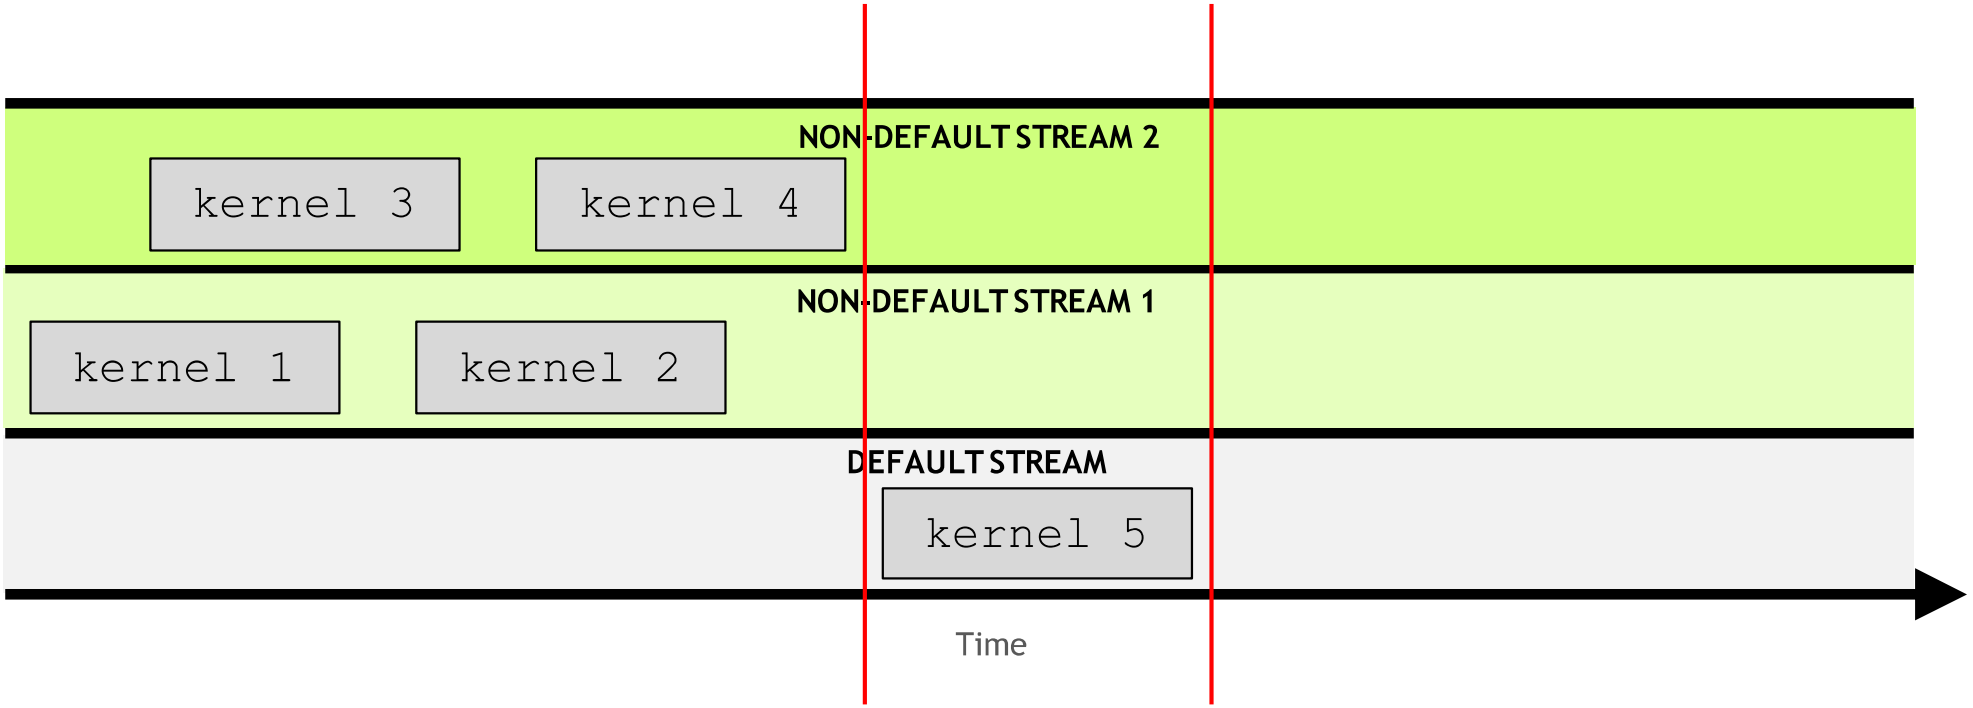

The default stream is special: **it blocks all kernels in all other streams**

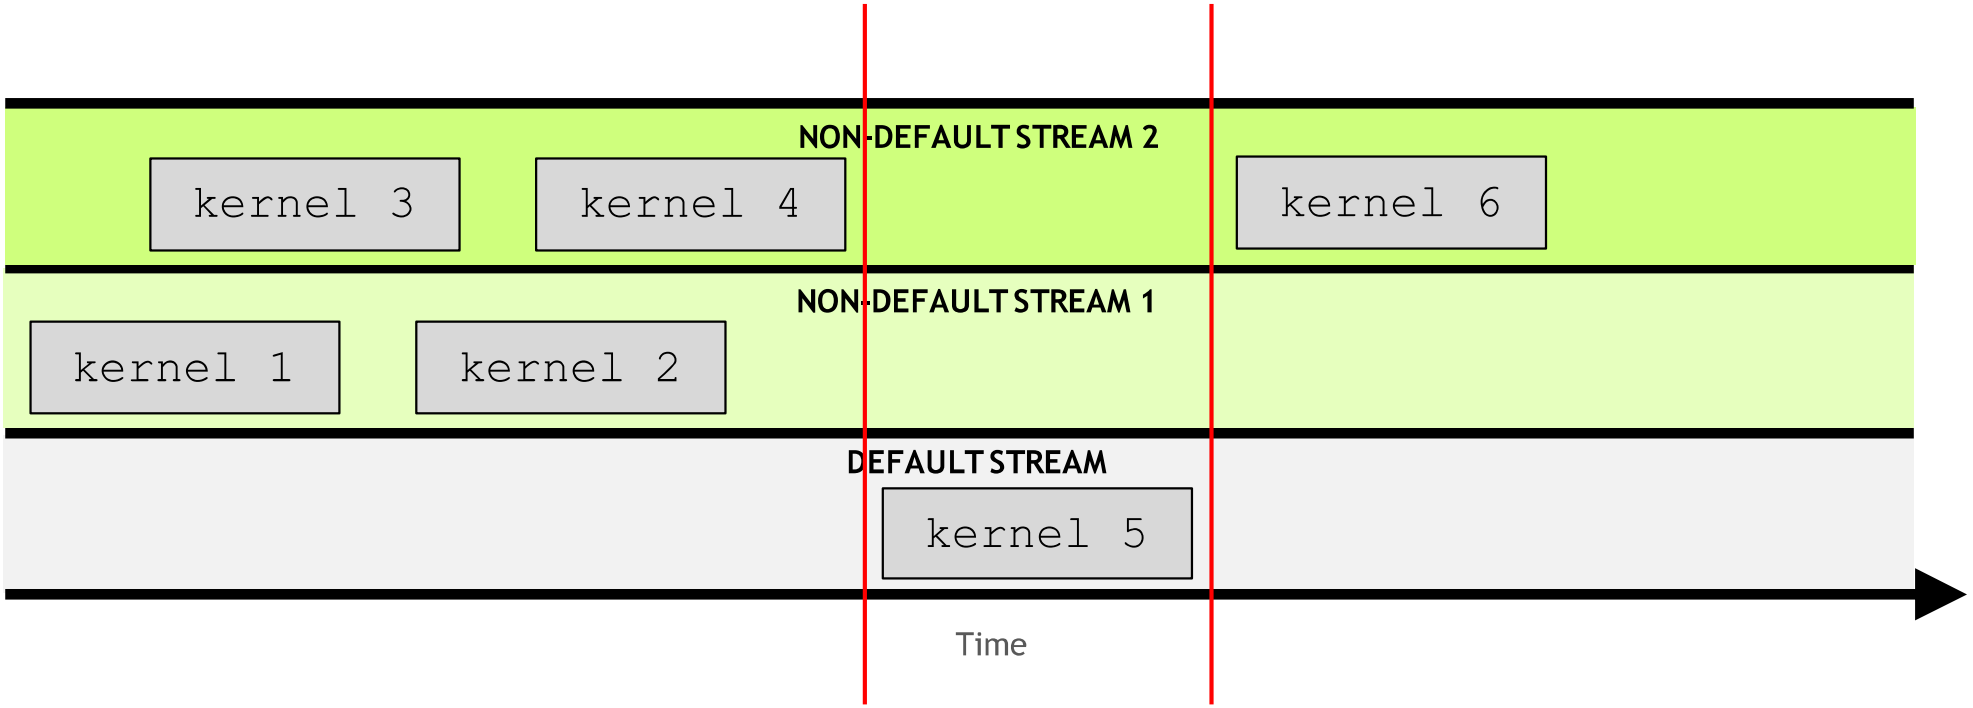

The default stream is special: **it blocks all kernels in all other streams**

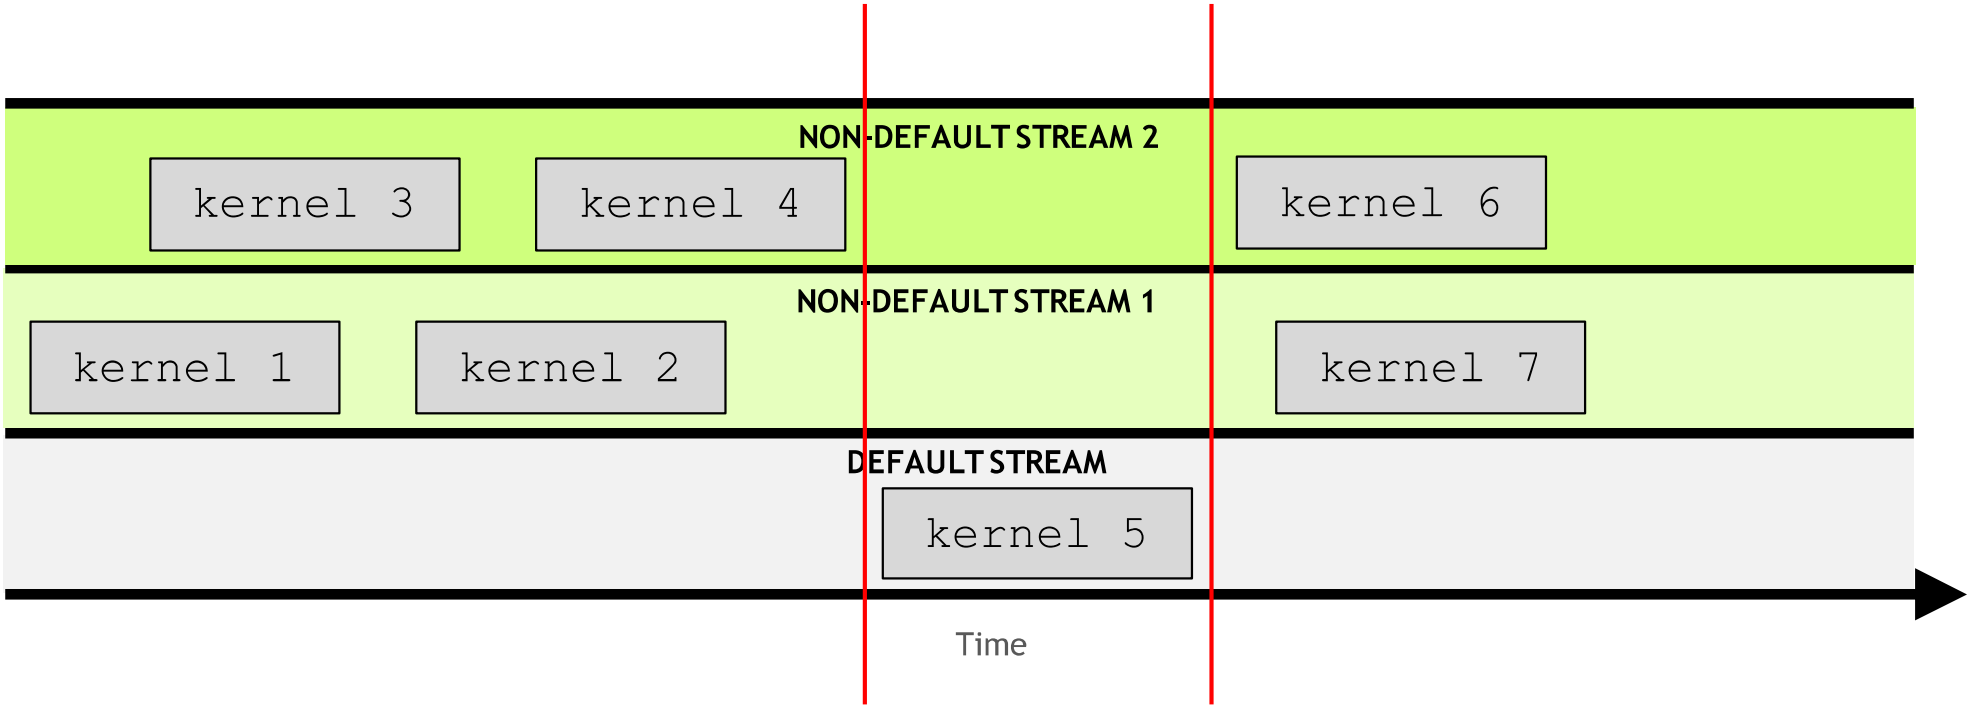

The default stream is special: **it blocks all kernels in all other streams**

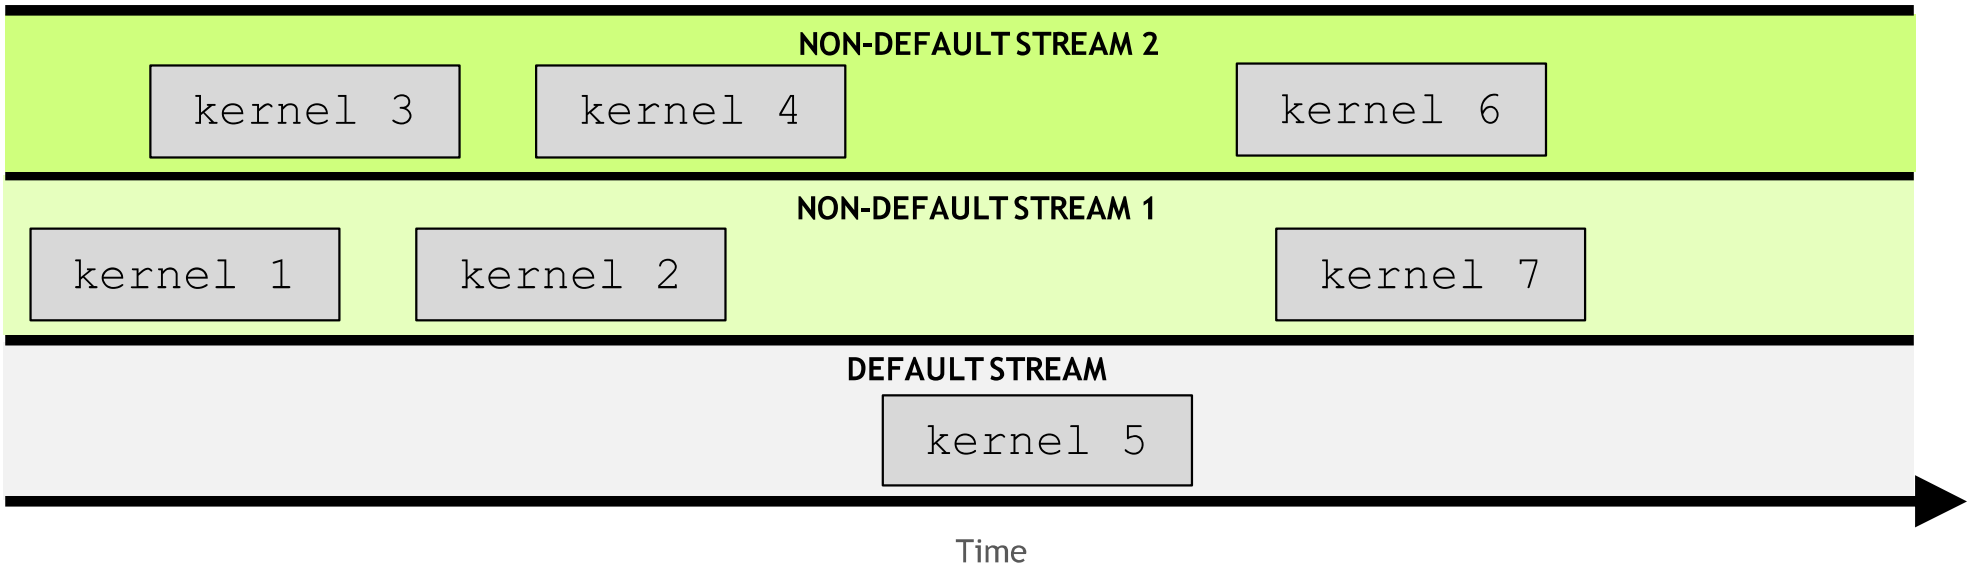

Non-Unified Memory

Memory can be allocated directly to the GPU with **`cudaMalloc`**

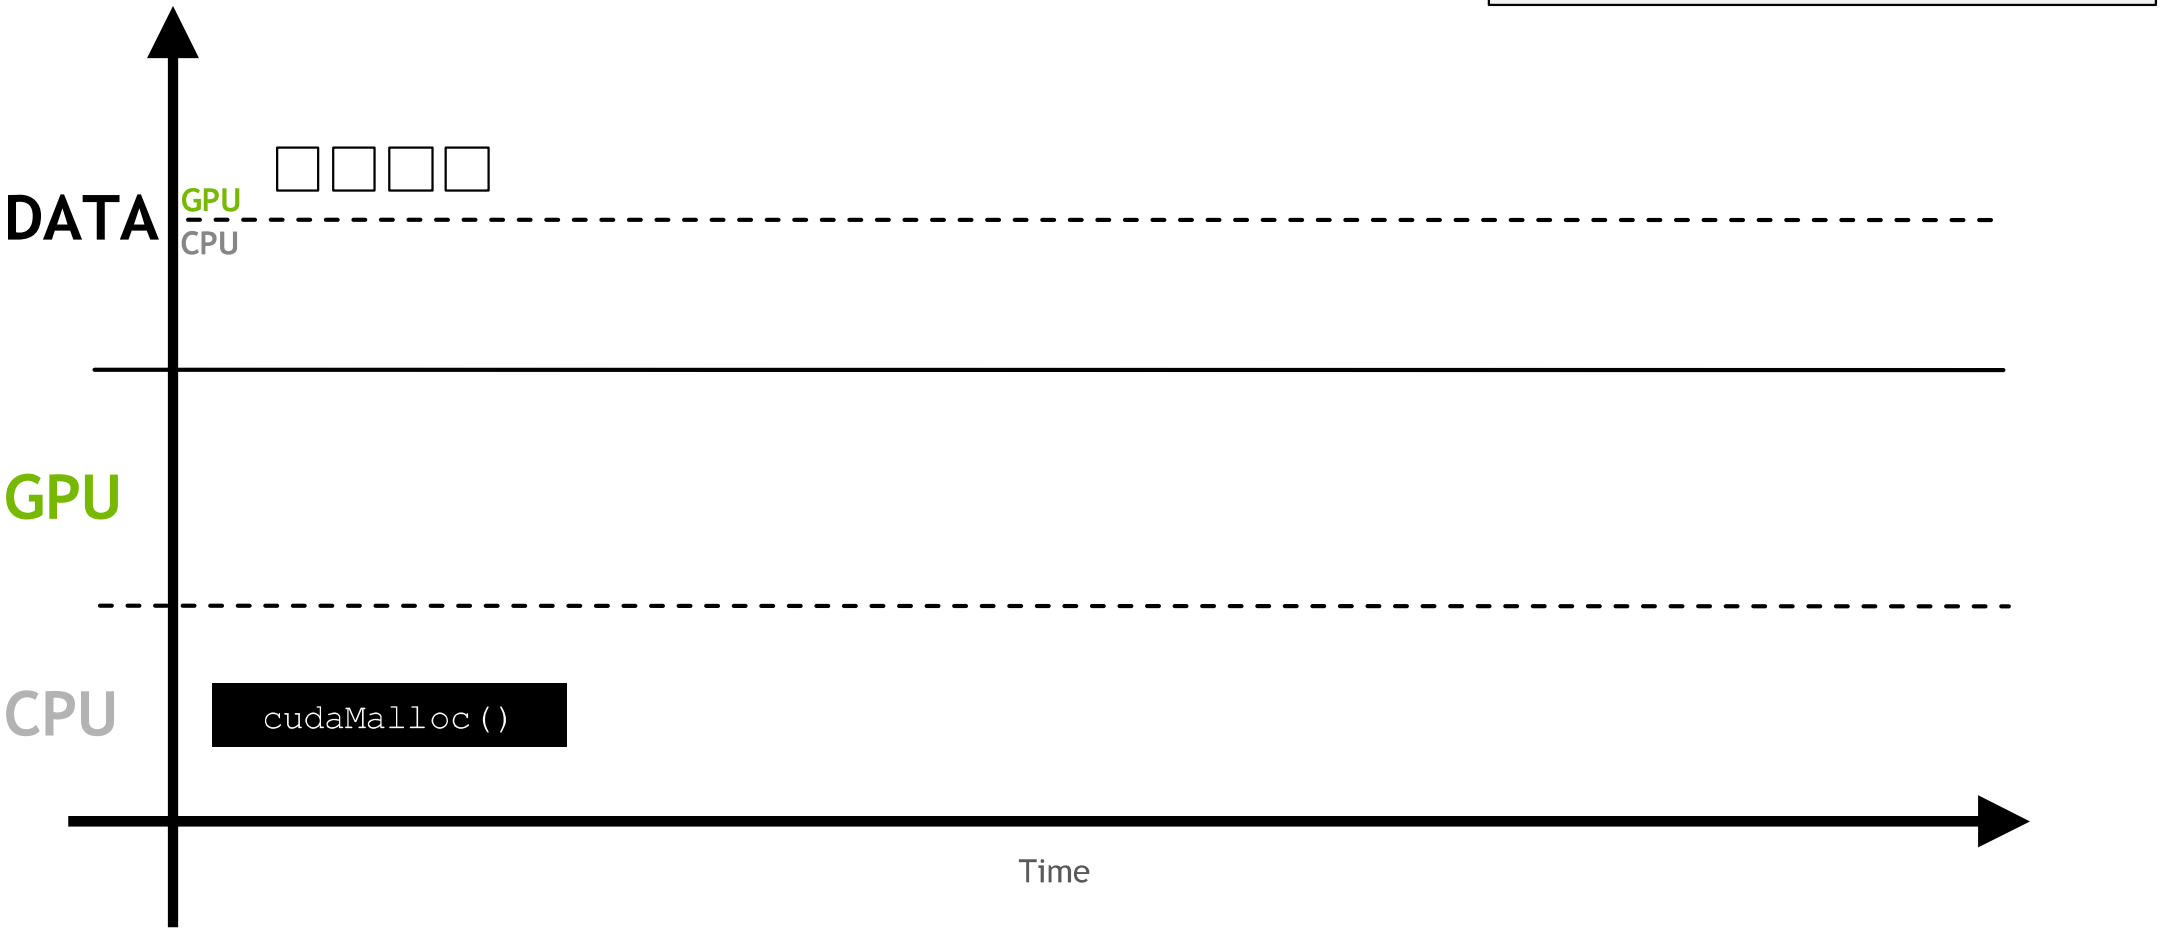

Memory can be allocated directly to the host with **`cudaMallocHost`**

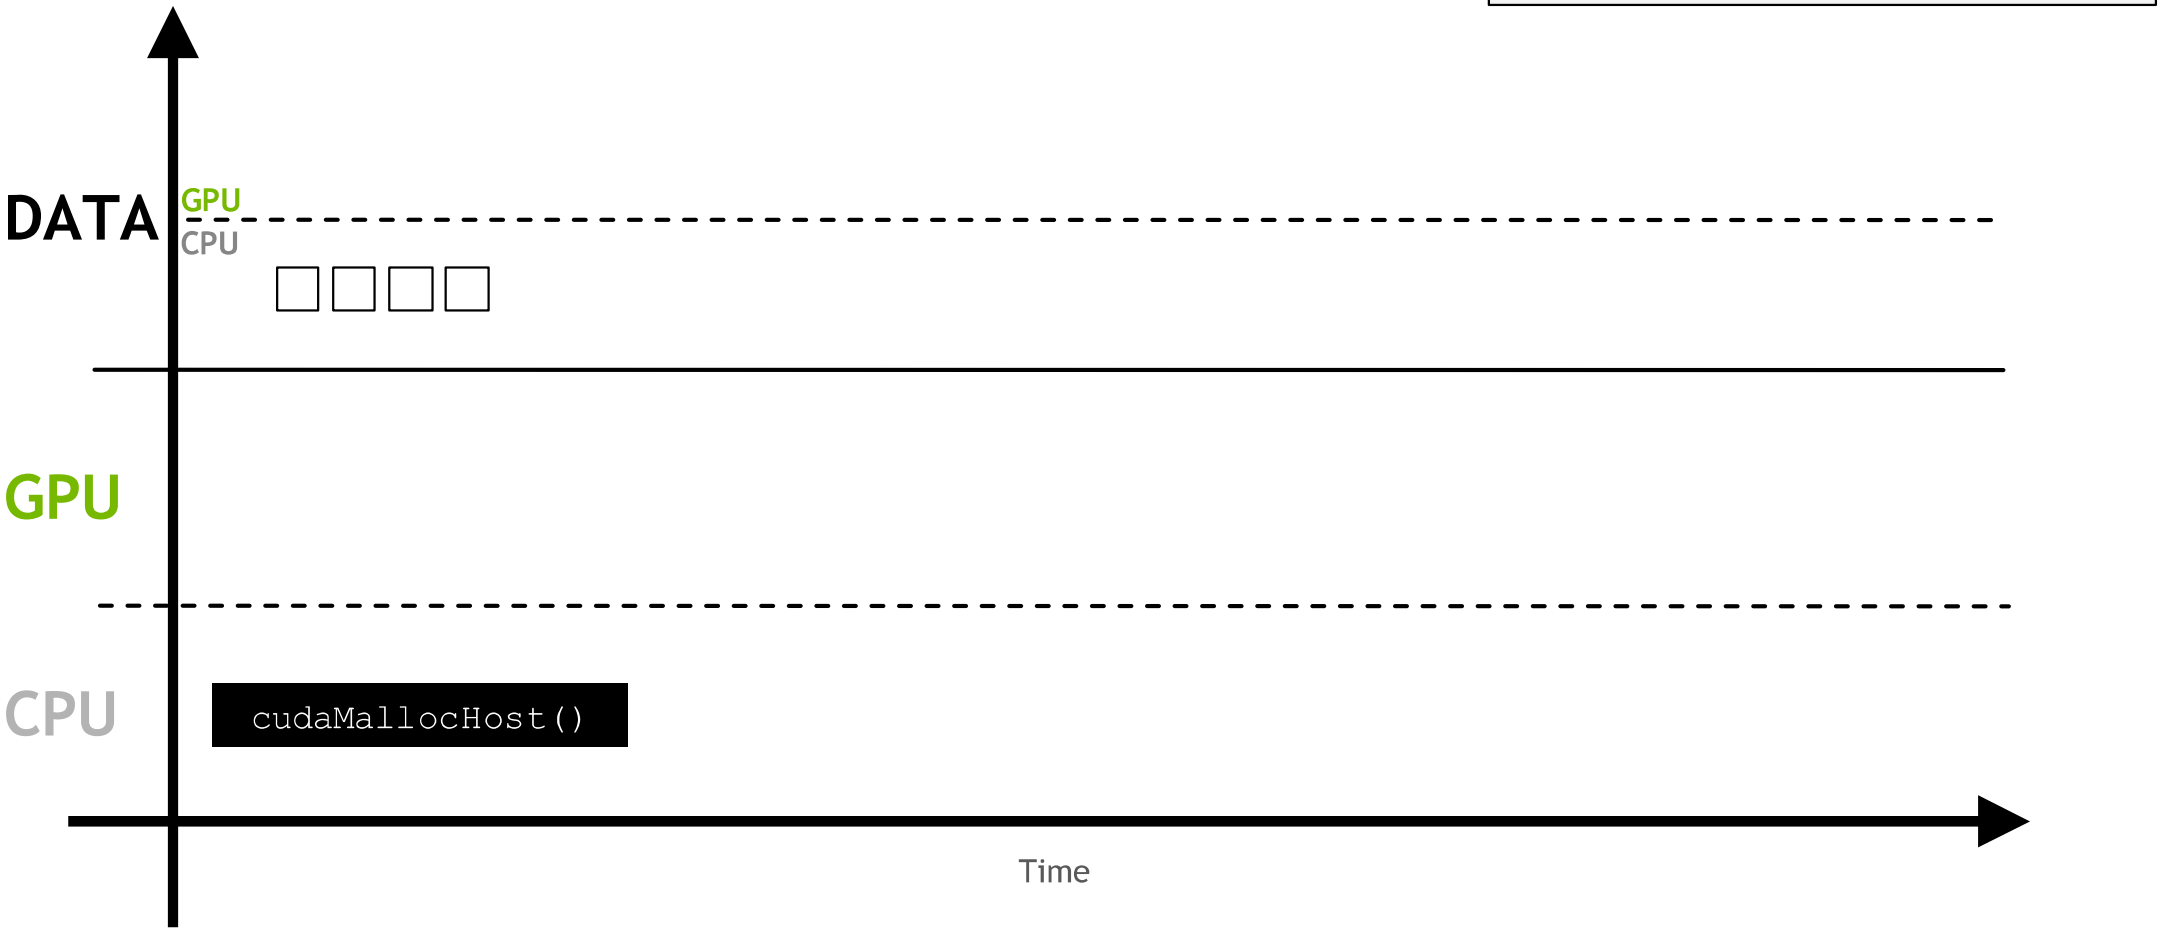

Memory allocated in either of these ways can be **copied** to other locations in the system with **`cudaMemCpy`**

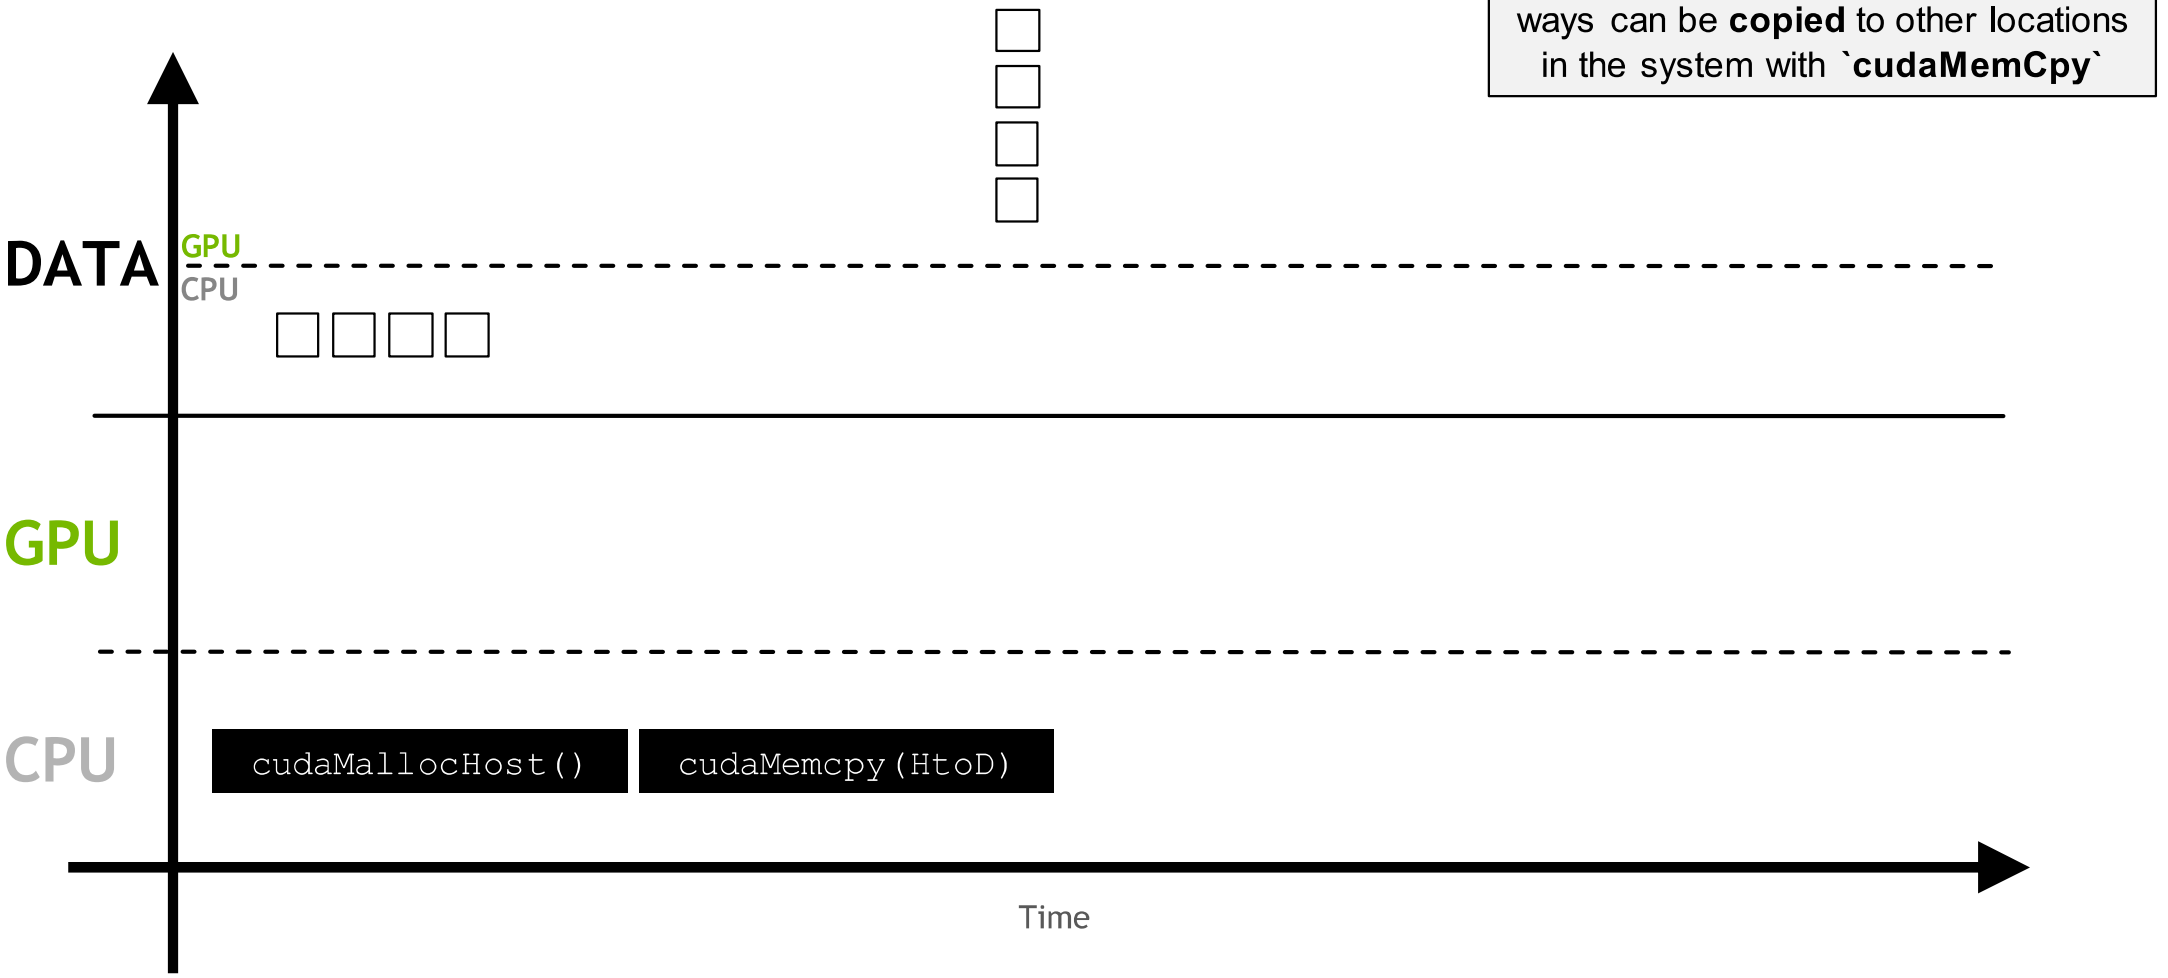

Copying leaves 2 copies in of in the system

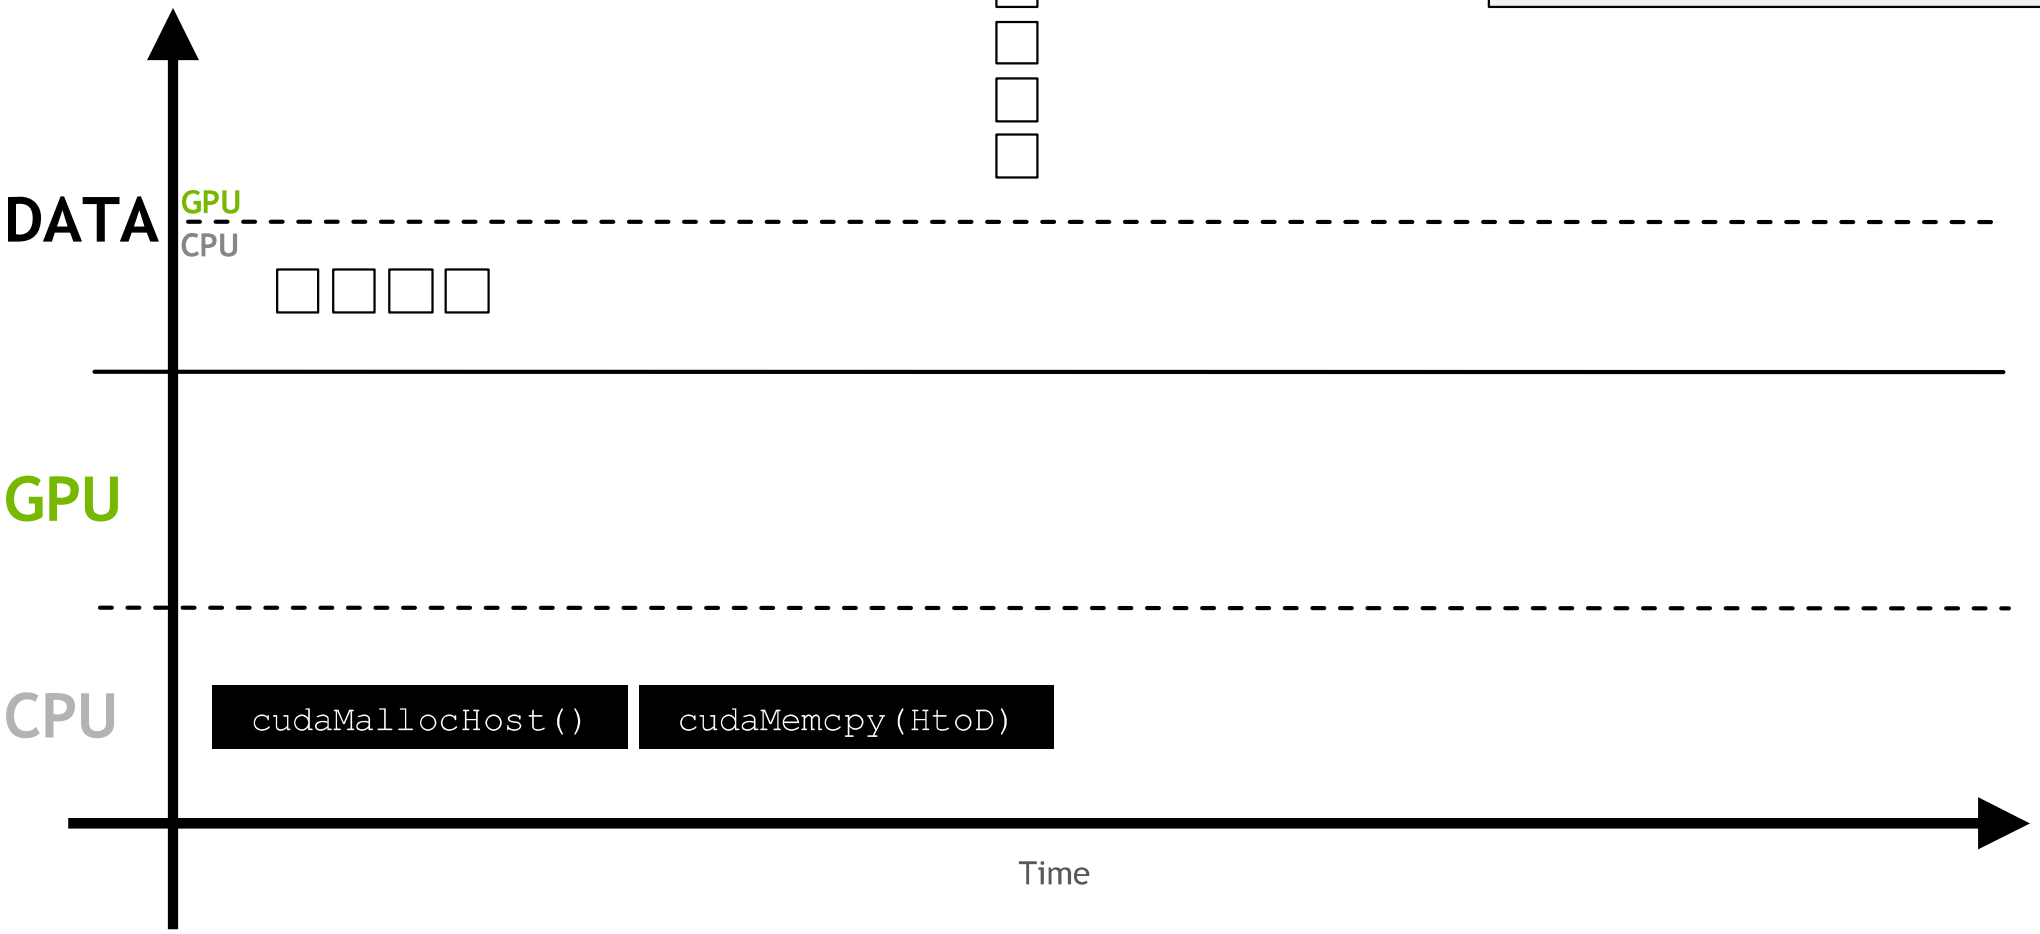

cudaMemcpyAsync

`cudaMemcpyAsync` can asynchronously transfer memory over a non-default stream

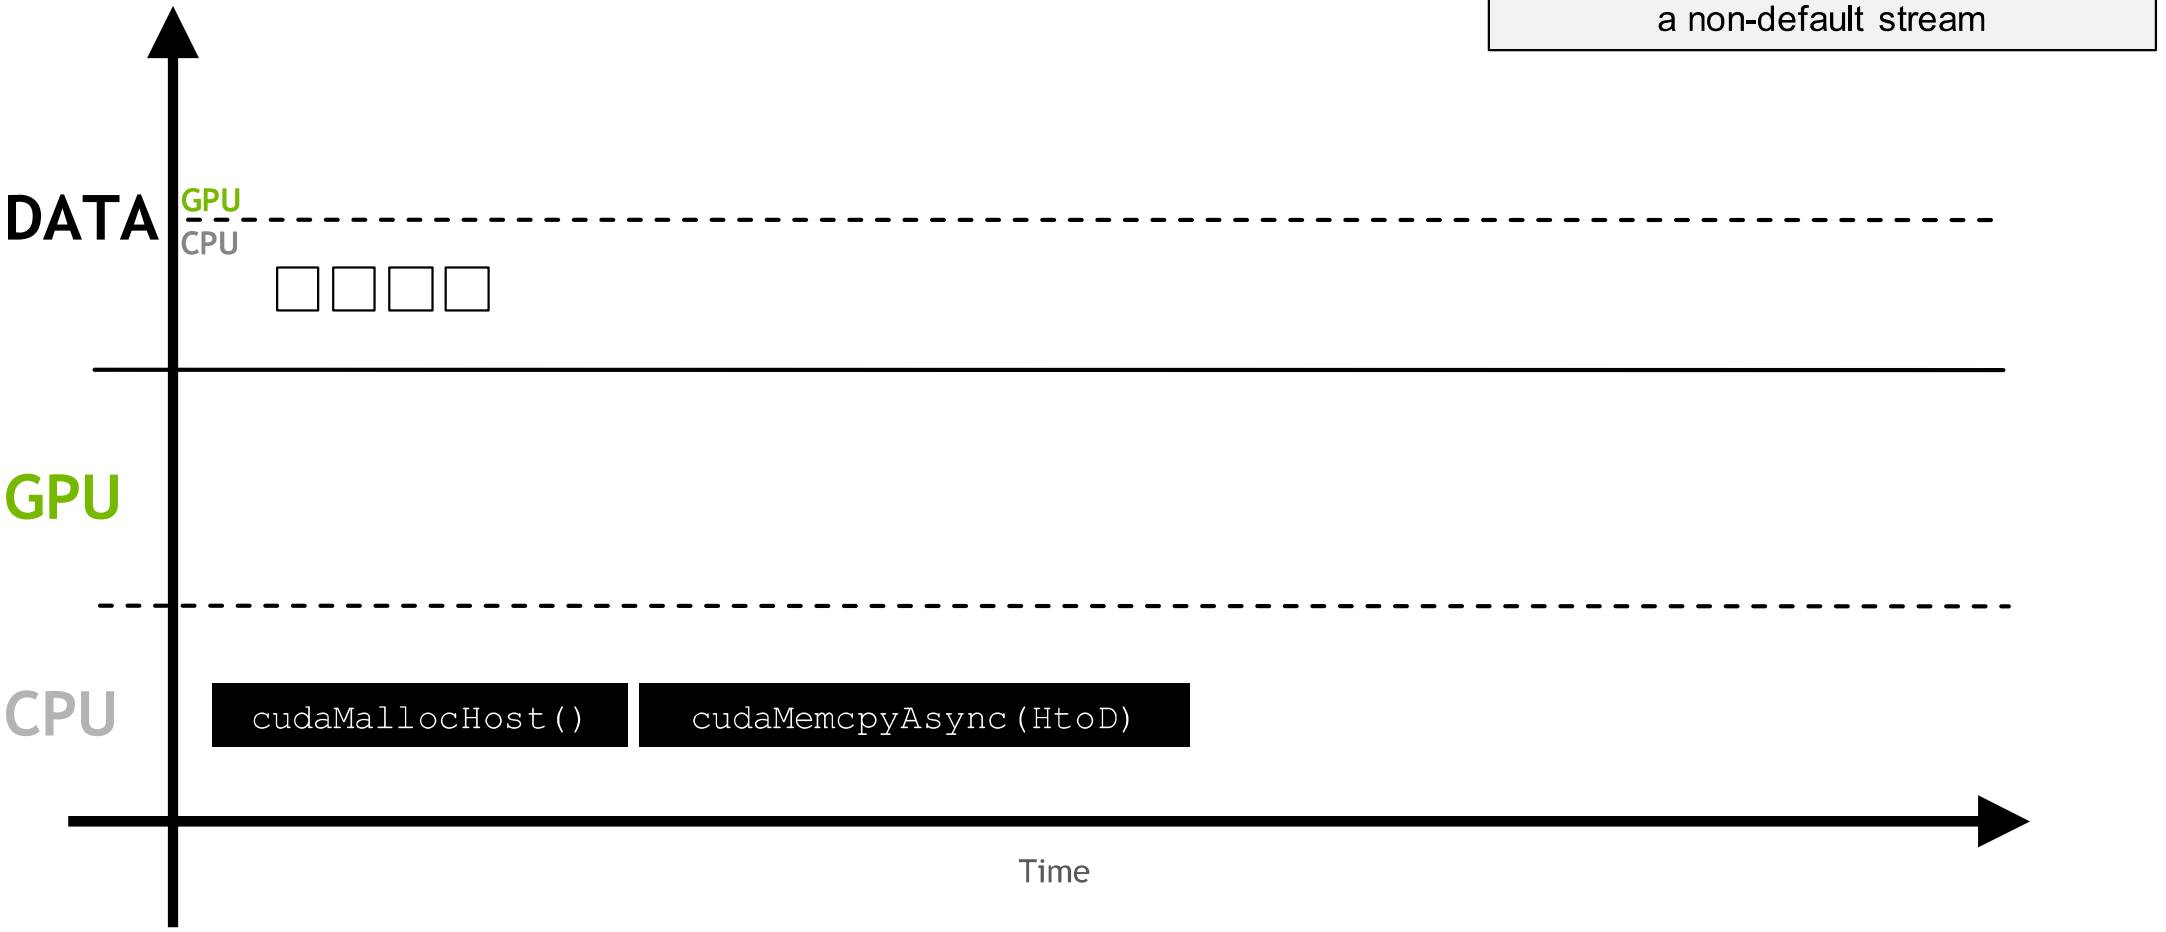

This can allow the **overlapping** memory copies and computation

This can allow the **overlapping** memory copies and computation

This can allow the **overlapping** memory copies and computation

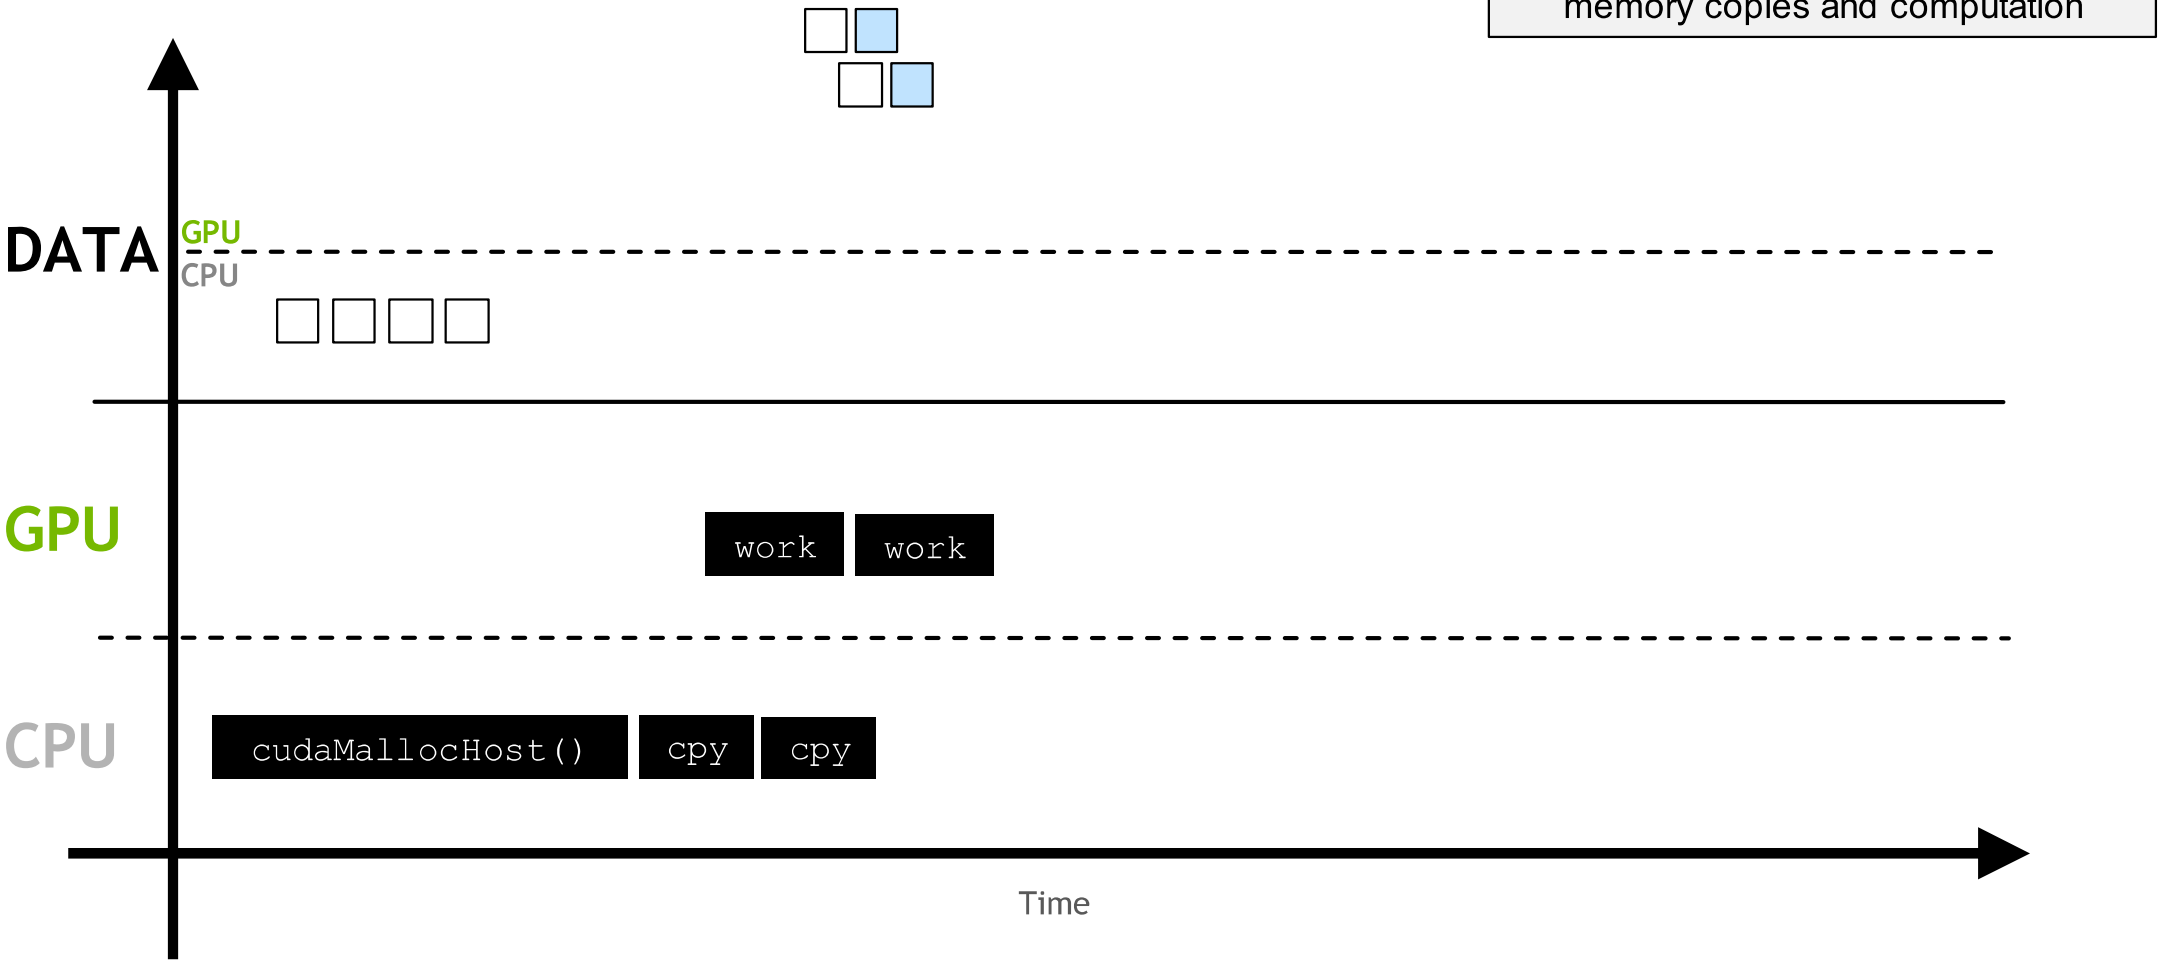

This can allow the **overlapping** memory copies and computation

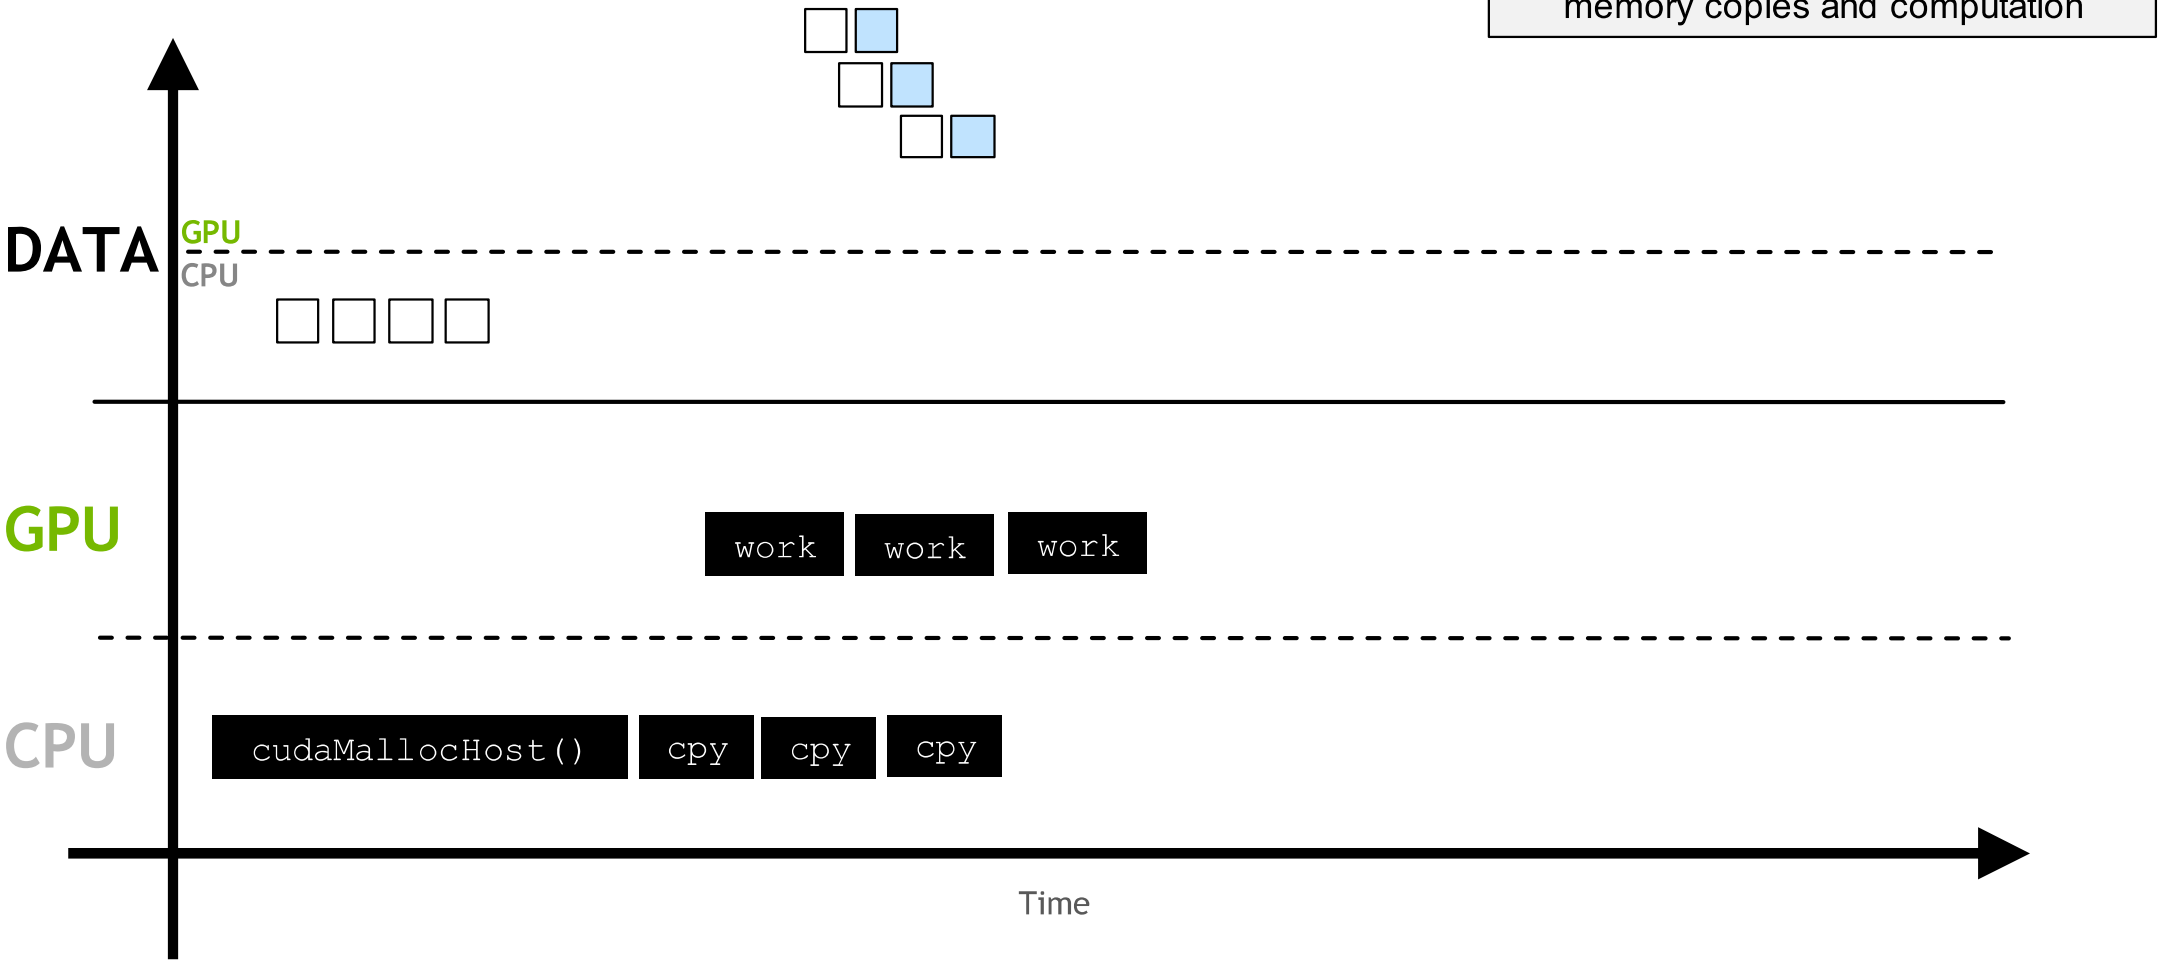

This can allow the **overlapping** memory copies and computation

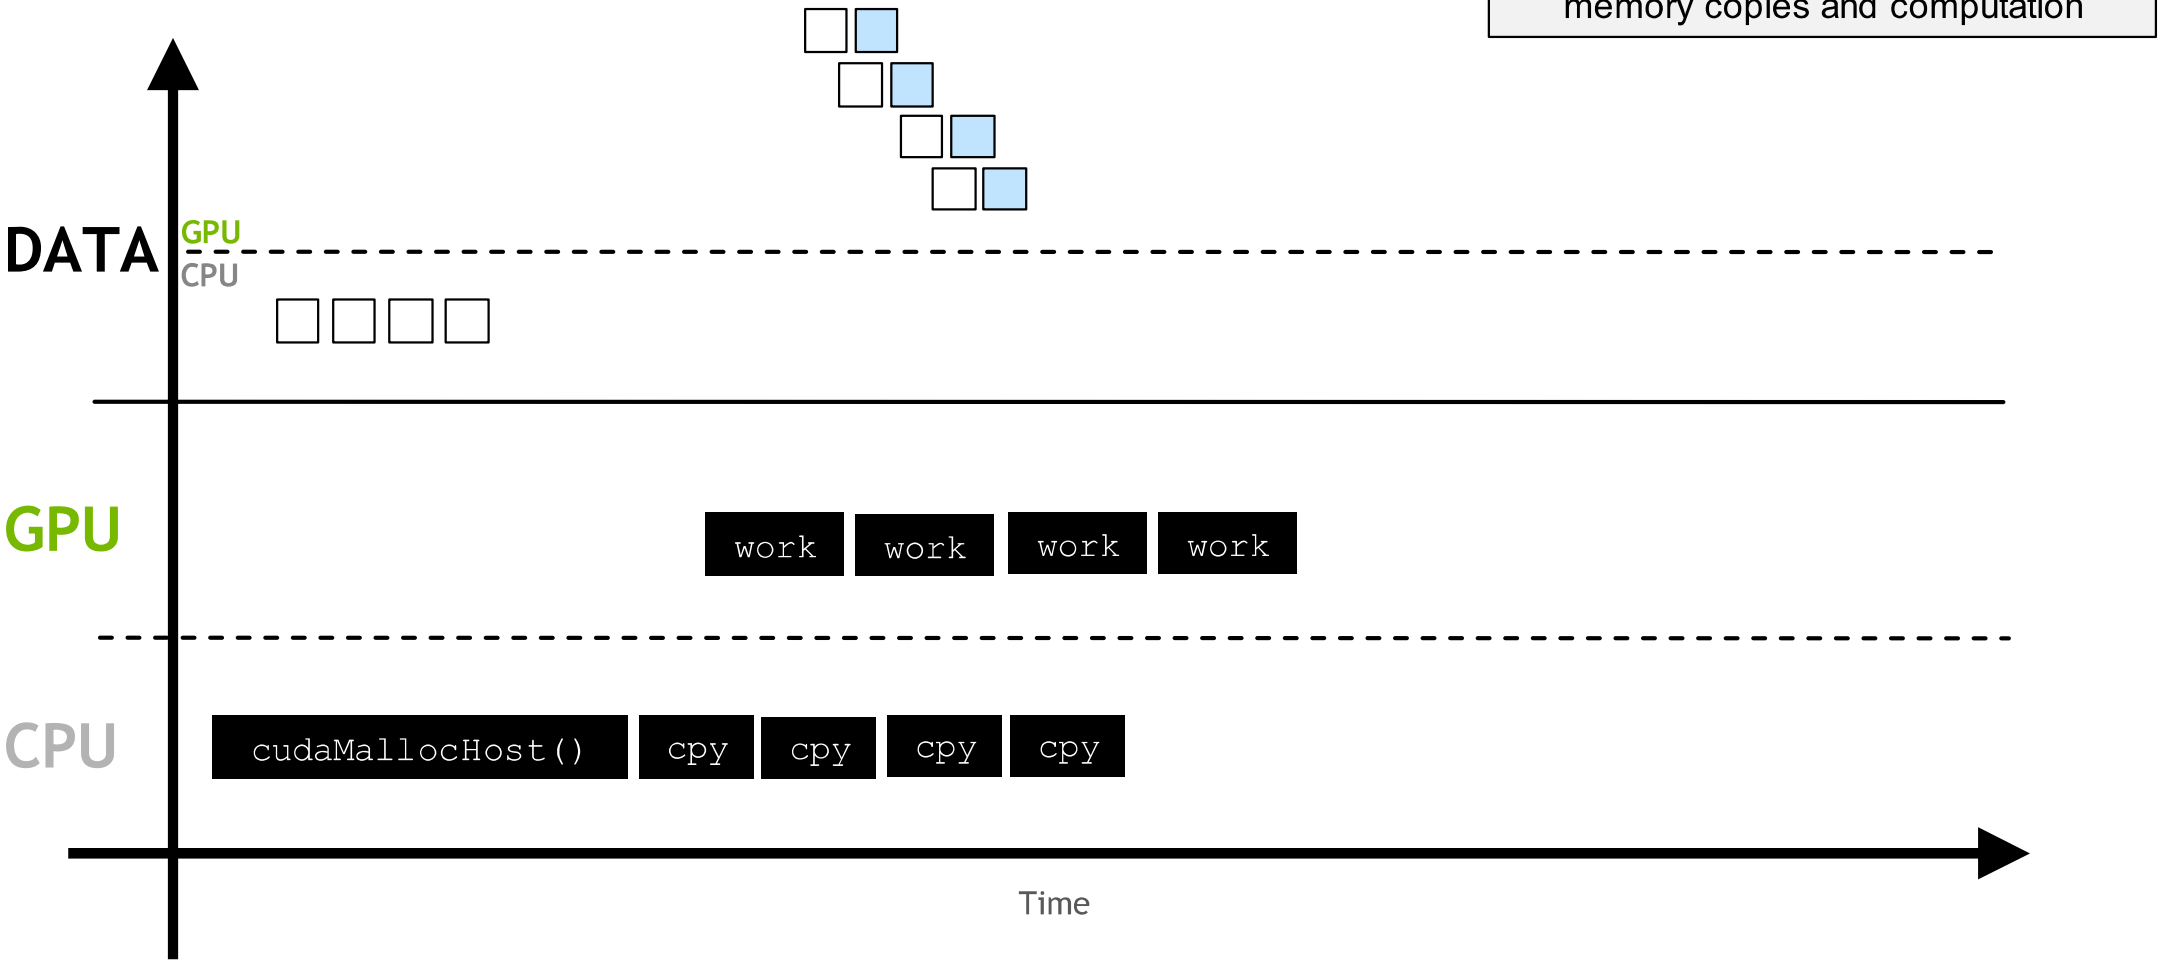

DEEP
LEARNING
INSTITUTE

www.nvidia.com/dli